

# ikEASY-Diabetes: System Documentation Suite

## Document Set Overview

This documentation suite provides a complete technical and clinical foundation for developing EASY-Diabetes, an advanced AI-powered clinical decision support system for comprehensive diabetes management. The following documents are designed to guide development teams through implementation while ensuring clinical accuracy, technical excellence, and regulatory compliance.

---

## Document 1: Executive Summary and Project Overview

### Project Vision

EASY-Diabetes represents the next generation of clinical decision support systems, leveraging cutting-edge AI technologies, real-time evidence synthesis, and personalized medicine approaches to revolutionize diabetes care. The system aims to reduce the global burden of diabetes by providing clinicians with intelligent, evidence-based recommendations while empowering patients with personalized insights and predictive analytics.

### Core Objectives

1. **Clinical Excellence:** Achieve >95% guideline concordance while personalizing care based on individual patient characteristics
2. **Predictive Analytics:** Provide 2-4 hour glucose predictions with >90% accuracy
3. **Treatment Optimization:** Reduce HbA1c by average 1.0-1.5% while minimizing hypoglycemia
4. **Scalability:** Support 1M+ concurrent users across multiple healthcare systems
5. **Interoperability:** Seamless integration with 95% of major EHR systems
6. **Regulatory Compliance:** FDA Class II clearance, CE marking, and full HIPAA/GDPR compliance

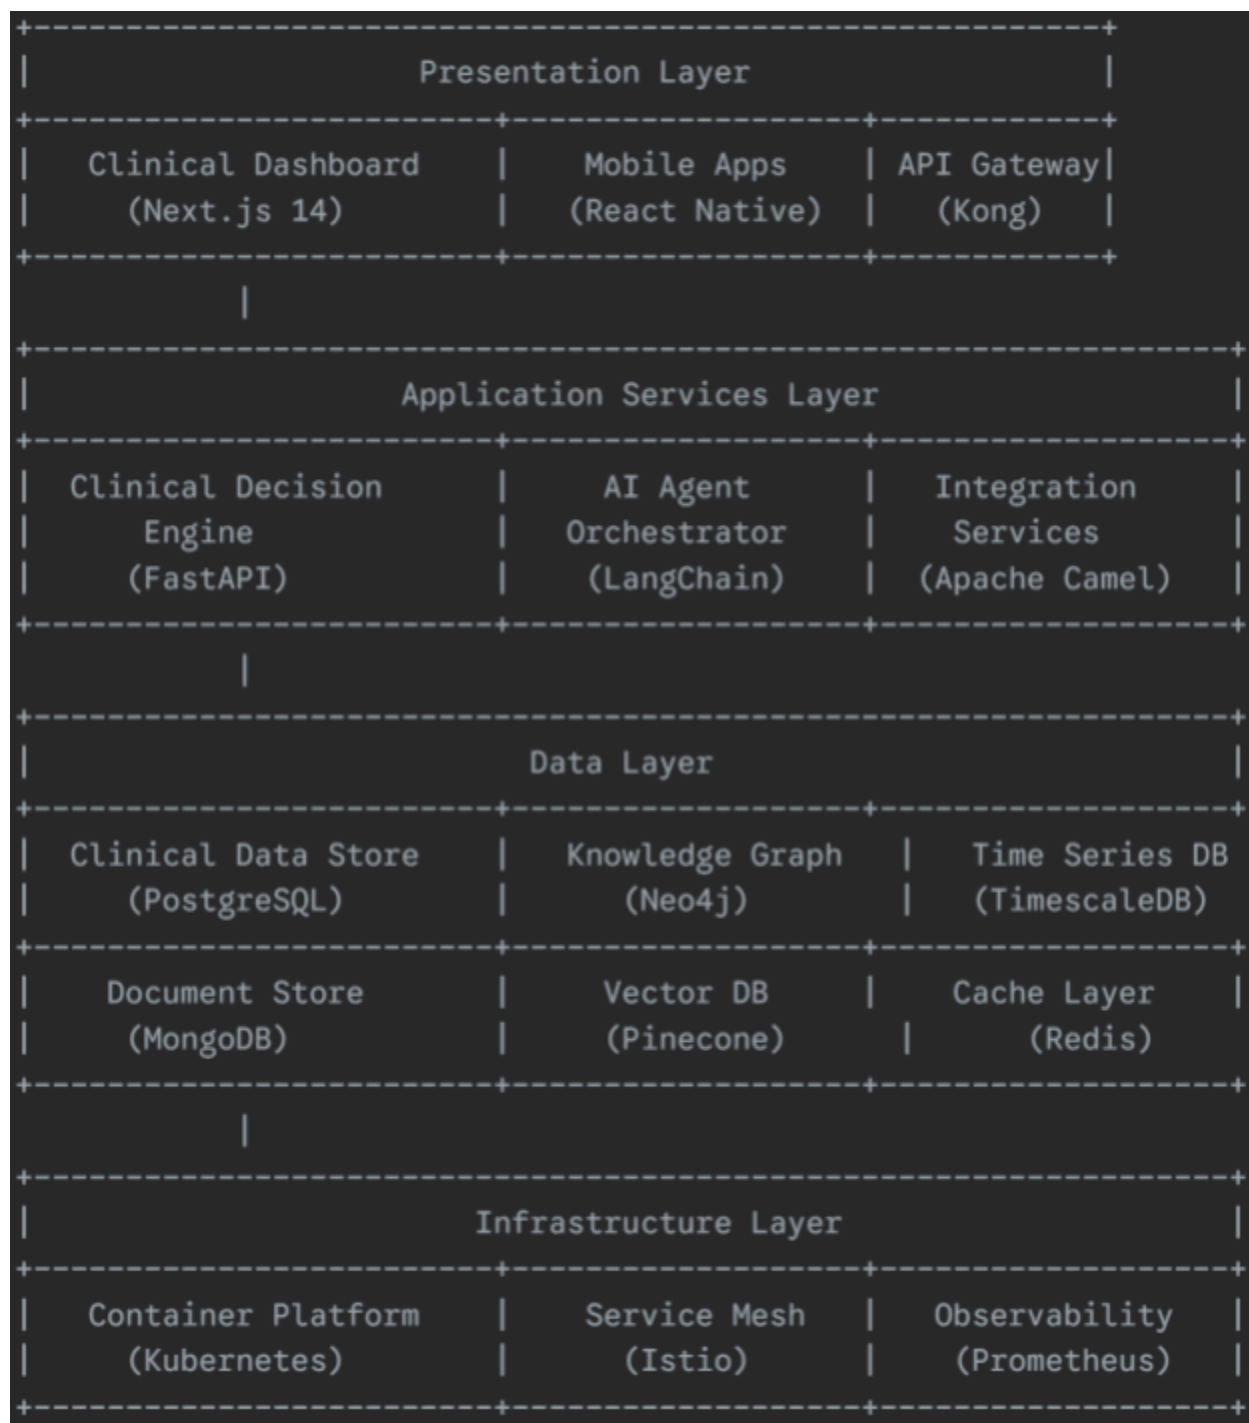
## Key Differentiators

- Multi-agent AI architecture with specialized clinical reasoning
  - Real-time integration of latest clinical evidence
  - Federated learning for continuous improvement
  - Explainable AI with full audit trails
  - Predictive analytics for proactive care
  - Comprehensive safety validation system
- 

# Document 2: System Architecture Document

## 1. High-Level Architecture Overview

### 1.1 System Components



## 1.2 Microservices Architecture

### Core Services

#### Clinical Decision Service

- **Technology:** Python 3.11, FastAPI, Pydantic v2

- **Responsibilities:** Clinical logic, guideline processing, recommendation generation
- **Scaling:** Horizontal, 3-20 instances based on load
- **Dependencies:** Knowledge Graph, Clinical Data Store, AI Services

### AI Agent Orchestrator Service

- **Technology:** Python 3.11, LangChain, Ray Serve
- **Responsibilities:** Multi-agent coordination, reasoning chains, explanation generation
- **Scaling:** GPU-enabled pods, 2-10 instances
- **Dependencies:** Vector DB, ML Model Registry, Clinical Decision Service

### Patient Data Service

- **Technology:** Go 1.21, gRPC, Protocol Buffers
- **Responsibilities:** Patient data CRUD, data validation, privacy enforcement
- **Scaling:** Horizontal, 5-50 instances
- **Dependencies:** PostgreSQL, Redis, Audit Service

### Integration Service

- **Technology:** Java 17, Apache Camel, Spring Boot
- **Responsibilities:** EHR integration, HL7/FHIR processing, external API management
- **Scaling:** Horizontal, 3-15 instances
- **Dependencies:** Message Queue, Transformation Engine

### Prediction Service

- **Technology:** Python 3.11, PyTorch, ONNX Runtime
- **Responsibilities:** Glucose prediction, risk scoring, outcome forecasting
- **Scaling:** GPU-enabled, 2-8 instances
- **Dependencies:** Time Series DB, Feature Store, Model Registry

### Monitoring Service

- **Technology:** Rust 1.75, Actix-web
- **Responsibilities:** Real-time monitoring, alert generation, anomaly detection
- **Scaling:** Horizontal, 5-20 instances
- **Dependencies:** Time Series DB, Notification Service

## 1.3 Data Architecture

### Primary Databases

#### PostgreSQL 16 (Clinical Data)

-- Core schema design

```

CREATE SCHEMA clinical;
CREATE SCHEMA patient;
CREATE SCHEMA audit;

-- Enable extensions
CREATE EXTENSION IF NOT EXISTS "pgcrypto";
CREATE EXTENSION IF NOT EXISTS "uuid-ossf";
CREATE EXTENSION IF NOT EXISTS "btree_gin";
CREATE EXTENSION IF NOT EXISTS "pg_trgm";

-- Partitioning strategy
CREATE TABLE patient.clinical_measurements (
  id UUID DEFAULT uuid_generate_v4(),
  patient_id UUID NOT NULL,
  measurement_type VARCHAR(50) NOT NULL,
  value NUMERIC(10,3) NOT NULL,
  unit VARCHAR(20) NOT NULL,
  measured_at TIMESTAMPTZ NOT NULL,
  metadata JSONB,
  created_at TIMESTAMPTZ DEFAULT NOW()
) PARTITION BY RANGE (measured_at);

-- Create monthly partitions
CREATE TABLE patient.clinical_measurements_2024_01
  PARTITION OF patient.clinical_measurements
  FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');

```

## Neo4j 5.0 (Knowledge Graph)

```

// Clinical knowledge graph schema
CREATE CONSTRAINT patient_id ON (p:Patient) ASSERT p.id IS UNIQUE;
CREATE CONSTRAINT medication_id ON (m:Medication) ASSERT m.id IS UNIQUE;
CREATE CONSTRAINT condition_id ON (c:Condition) ASSERT c.id IS UNIQUE;

// Relationship types
CREATE INDEX ON :TAKES_MEDICATION(start_date);
CREATE INDEX ON :HAS_CONDITION(diagnosed_date);
CREATE INDEX ON :SIMILAR_TO(similarity_score);

```

## TimescaleDB (Time Series)

```

-- Hypertable for continuous data
CREATE TABLE sensor_data (

```

```

time TIMESTAMPTZ NOT NULL,
device_id UUID NOT NULL,
patient_id UUID NOT NULL,
metric_type VARCHAR(50) NOT NULL,
value DOUBLE PRECISION NOT NULL,
quality_flag INTEGER
);

SELECT create_hypertable('sensor_data', 'time');

-- Compression policy
ALTER TABLE sensor_data SET (
    timescaledb.compress,
    timescaledb.compress_segmentby = 'device_id, patient_id'
);

```

## 1.4 AI/ML Architecture

### Multi-Agent System Design

#### Agent Hierarchy

```

class AgentHierarchy:
    master_orchestrator = {
        "model": "gpt-4-turbo-128k",
        "temperature": 0.1,
        "role": "Clinical decision coordination"
    }

    specialist_agents = {
        "clinical_assessment": {
            "model": "claude-4-opus",
            "tools": ["lab_analyzer", "symptom_checker", "risk_calculator"],
            "specialization": "Patient evaluation and diagnosis support"
        },
        "treatment_optimization": {
            "model": "med-palm-2",
            "tools": ["drug_database", "interaction_checker", "dosing_calculator"],
            "specialization": "Medication selection and optimization"
        },
        "guideline_compliance": {
            "model": "gpt-4-medical-fine-tuned",
            "tools": ["guideline_search", "evidence_ranker", "contradiction_detector"],
            "specialization": "Ensuring evidence-based recommendations"
        }
    }

```

```

    },
    "safety_validation": {
      "model": "claude-4-safety",
      "tools": ["contraindication_checker", "allergy_validator", "risk_assessor"],
      "specialization": "Safety verification and risk mitigation"
    }
  }
}

```

## ML Model Pipeline

### Training Infrastructure

- **Platform:** Kubeflow on Kubernetes
- **Experiment Tracking:** MLflow + DVC
- **Model Registry:** MLflow Model Registry
- **Feature Store:** Feast
- **Distributed Training:** Ray + Horovod

### Model Deployment

- **Serving:** Ray Serve + TorchServe
- **A/B Testing:** Seldon Core
- **Edge Deployment:** ONNX Runtime + TensorFlow Lite
- **Model Monitoring:** Evidently AI + Grafana

## 1.5 Security Architecture

### Zero Trust Security Model

#### Network Security

```

# Istio service mesh configuration
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
spec:
  mtls:
    mode: STRICT
---
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: clinical-services

```

spec:  
 selector:  
 matchLabels:  
 app: clinical-decision-service  
 rules:  
 - from:  
 - source:  
 principals: ["cluster.local/ns/easy-diabetes/sa/api-gateway"]  
 to:  
 - operation:  
 methods: ["GET", "POST"]

## Data Encryption

- **At Rest:** AES-256-GCM with AWS KMS/Azure Key Vault
- **In Transit:** TLS 1.3 with mutual authentication
- **Application Level:** Field-level encryption for PII
- **Key Rotation:** Automated 90-day rotation

## Access Control

- **Authentication:** OAuth 2.0 + OpenID Connect
- **Authorization:** RBAC with ABAC policies
- **MFA:** TOTP/WebAuthn for clinical users
- **Session Management:** Redis-backed with 15-minute idle timeout

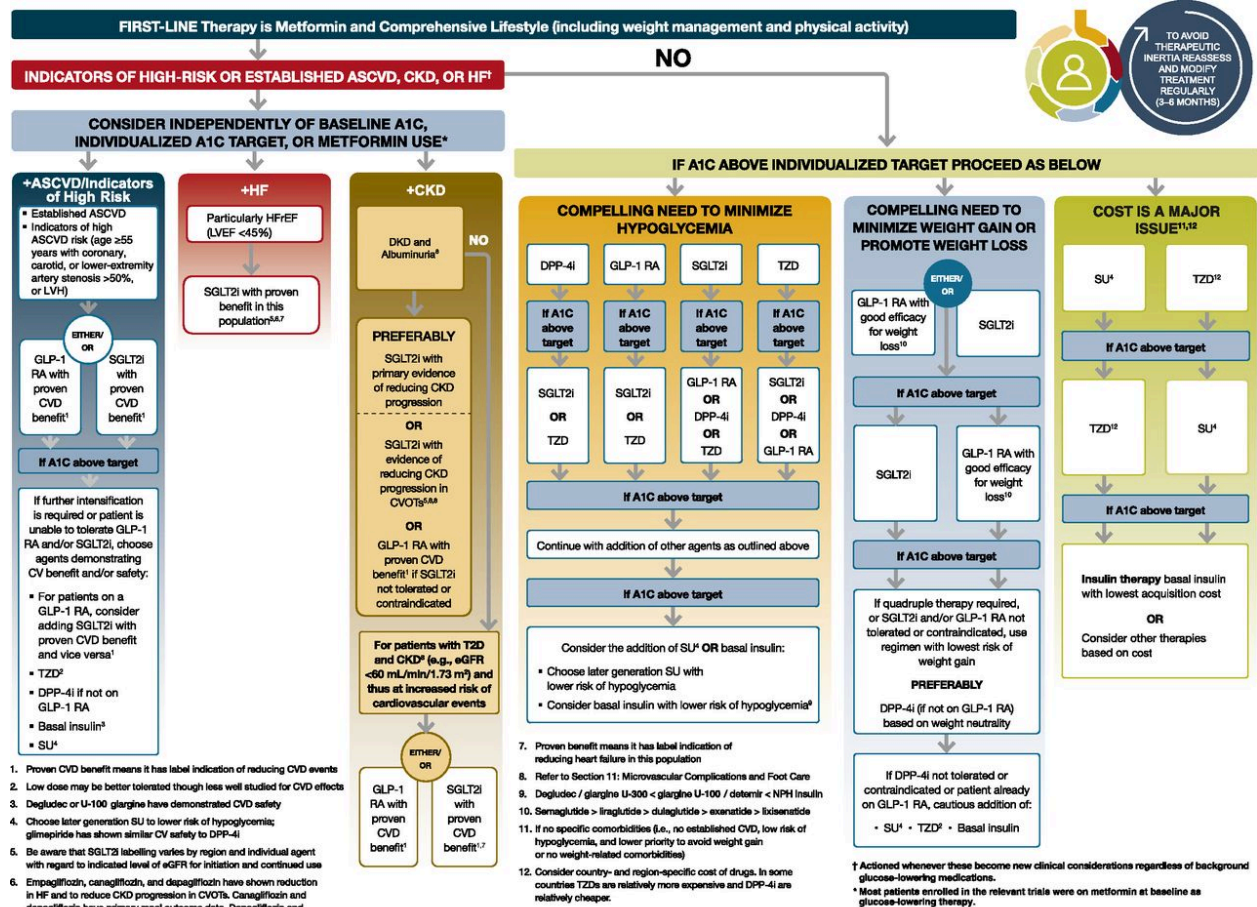
---

# Document 3: Medical Knowledge Base and Clinical Guidelines

## 3.1 Comprehensive Clinical Guidelines Integration

### Type 2 Diabetes Management Guidelines





## American Diabetes Association (ADA) Standards of Care 2024

- **Glycemic Targets:**
  - General adult target: HbA1c <7.0% (53 mmol/mol)
  - Individualized targets based on:
    - Life expectancy
    - Comorbidities
    - Hypoglycemia risk
    - Disease duration
    - Patient preferences
  - Time in Range (TIR) target: >70% (70-180 mg/dL)
  - Time below range: <4% (<70 mg/dL), <1% (<54 mg/dL)

## Medication Algorithm (2024 Update):

class ADAMedicationAlgorithm:

```

first_line = {
    "medication": "Metformin",
    "contraindications": ["eGFR <30", "acidosis_risk", "hypoxia"],
    "starting_dose": "500mg daily or BID",

```

```

    "titration": "Increase by 500mg weekly to max 2000mg"
  }

  second_line_options = {
    "ASCVD_or_high_risk": ["GLP-1_RA", "SGLT2i"],
    "HF_or_CKD": ["SGLT2i"],
    "weight_management_priority": ["GLP-1_RA", "dual_GIP_GLP1"],
    "hypoglycemia_concern": ["DPP-4i", "GLP-1_RA", "SGLT2i", "TZD"],
    "cost_priority": ["Sulfonylurea", "TZD"]
  }

  injectable_therapies = {
    "GLP-1_before_insulin": True,
    "basal_insulin_start": "10_units_or_0.1-0.2_units/kg",
    "prandial_addition": "4_units_or_10%_basal_dose"
  }

```

## EASD/ADA Consensus Report 2023

- Emphasis on cardiorenal protection
- Early combination therapy for HbA1c >1.5% above target
- Continuous glucose monitoring for all on intensive insulin
- Annual screening updates:
  - Retinopathy: Starting at diagnosis
  - Nephropathy: Annual eGFR and UACR
  - Neuropathy: 10-g monofilament + one additional test
  - ASCVD: Lipids annually, more frequent if abnormal

## Type 1 Diabetes Management

### ADA/EASD Type 1 Consensus 2024

- **Insulin Regimens:**
  - Multiple daily injections (MDI): ≥4 injections/day
  - Continuous subcutaneous insulin infusion (CSII)
  - Automated insulin delivery (AID) systems preferred

### Insulin Dosing Algorithms:

```

class Type1InsulinAlgorithms:
    total_daily_dose = {
        "initial": "0.4-0.5 units/kg/day",
        "honeymoon": "0.2-0.4 units/kg/day",
        "established": "0.5-1.0 units/kg/day",
    }

```

```

    "insulin_resistant": ">1.0 units/kg/day"
  }

  basal_bolus_split = {
    "basal": "40-50% of TDD",
    "bolus": "50-60% of TDD"
  }

  carb_ratios = {
    "starting": "1:15 (1 unit per 15g carbs)",
    "calculation": "450-500 / TDD",
    "adjustment": "10-20% changes based on patterns"
  }

  correction_factors = {
    "starting": "1:50 (1 unit drops BG 50 mg/dL)",
    "calculation": "1700-1800 / TDD",
    "timing": "Give with meal if BG >140 mg/dL"
  }

```

## Special Populations

### Pediatric Diabetes (ISPAD 2024)

- **Glycemic Targets:**
  - HbA1c <7.0% without significant hypoglycemia
  - TIR >70%, TBR <4%
  - Sensor use for all children with T1D

### Gestational Diabetes (ACOG/ADA 2024)

- **Screening:**
  - Universal screening at 24-28 weeks
  - Early screening if risk factors
- **Targets:**
  - Fasting <95 mg/dL
  - 1-hour postprandial <140 mg/dL
  - 2-hour postprandial <120 mg/dL

### Elderly Diabetes Management (AGS/ADA)

- **Healthy elderly:** Standard targets
- **Complex/intermediate health:** HbA1c <8.0%
- **Very complex/poor health:** HbA1c <8.5%

## Diabetes Technology Integration

### Continuous Glucose Monitoring (CGM)

```
class CGMIntegration:
    devices = {
        "dexcom_g7": {
            "warm_up": "30 minutes",
            "sensor_life": "10 days",
            "accuracy": "9.1% MARD",
            "features": ["predictive_alerts", "smartphone_direct"]
        },
        "freestyle_libre_3": {
            "warm_up": "60 minutes",
            "sensor_life": "14 days",
            "accuracy": "9.2% MARD",
            "features": ["real_time_continuous", "optional_alarms"]
        },
        "medtronic_guardian_4": {
            "warm_up": "2 hours",
            "sensor_life": "7 days",
            "accuracy": "9.1% MARD",
            "features": ["smart_guard_integration", "calibration_free"]
        }
    }

    metrics = {
        "time_in_range": {"target": 70-180, "goal": ">70%"},
        "time_below_range": {
            "level_1": {"range": 54-69, "goal": "<4%"},
            "level_2": {"range": "<54", "goal": "<1%"}
        },
        "time_above_range": {
            "level_1": {"range": 181-250, "goal": "<25%"},
            "level_2": {"range": ">250", "goal": "<5%"}
        },
        "glycemic_variability": {"cv": "<36%"}
    }
```

## Medication Database

### Comprehensive Drug Information

#### Metformin

```

class MetforminProfile:
    mechanism = "Decreases hepatic glucose production, increases insulin sensitivity"

    dosing = {
        "immediate_release": {
            "starting": "500mg daily or BID",
            "maximum": "2550mg daily (850mg TID)",
            "with_meals": True
        },
        "extended_release": {
            "starting": "500-1000mg daily",
            "maximum": "2000mg daily",
            "timing": "with_evening_meal"
        }
    }

    contraindications = [
        "eGFR <30 mL/min/1.73m²",
        "Acute metabolic acidosis",
        "Decompensated heart failure",
        "Severe hepatic impairment"
    ]

    precautions = {
        "eGFR_30_45": "Do not initiate; discontinue if falls below 30",
        "contrast_procedures": "Hold 48 hours post-procedure",
        "surgery": "Hold day of surgery",
        "b12_deficiency": "Monitor annually"
    }

    side_effects = {
        "common": ["GI upset (30%)", "Diarrhea", "Nausea", "Metallic taste"],
        "serious": ["Lactic acidosis (rare)", "B12 deficiency"]
    }

```

## SGLT2 Inhibitors

```

class SGLT2InhibitorProfiles:
    medications = {
        "empagliflozin": {
            "doses": ["10mg", "25mg"],
            "cv_outcome_trial": "EMPA-REG",
            "cv_death_reduction": "38%",
            "renal_benefits": "EMPA-KIDNEY positive"
        }
    }

```

```

},
"dapagliflozin": {
  "doses": ["5mg", "10mg"],
  "cv_outcome_trial": "DECLARE-TIMI",
  "hf_reduction": "27%",
  "renal_trial": "DAPA-CKD positive"
},
"canagliflozin": {
  "doses": ["100mg", "300mg"],
  "cv_outcome_trial": "CANVAS",
  "renal_trial": "CREDENCE positive",
  "warnings": ["amputation_risk", "fracture_risk"]
},
"ertugliflozin": {
  "doses": ["5mg", "15mg"],
  "cv_outcome_trial": "VERTIS-CV",
  "cv_neutral": True
}
}

```

```

class_effects = {
  "benefits": [
    "HbA1c reduction: 0.5-1.0%",
    "Weight loss: 2-3 kg",
    "BP reduction: 3-5 mmHg",
    "Cardiovascular protection",
    "Renal protection"
  ],
  "risks": [
    "Genital mycotic infections: 5-10%",
    "Euglycemic DKA: rare but serious",
    "Volume depletion",
    "Acute kidney injury (initial)",
    "Fournier's gangrene (rare)"
  ]
}

```

## GLP-1 Receptor Agonists

class GLP1AgonistProfiles:

```

  medications = {
    "semaglutide": {
      "formulations": {
        "ozempic": {"doses": ["0.25mg", "0.5mg", "1mg", "2mg"], "frequency": "weekly"},

```

```

    "rybelsus": {"doses": ["3mg", "7mg", "14mg"], "frequency": "daily_oral"},
    "wegovy": {"doses": ["up to 2.4mg"], "indication": "weight_management"}
  },
  "cv_trial": "SUSTAIN-6, PIONEER-6",
  "weight_loss": "5-15%"
},
"dulaglutide": {
  "doses": ["0.75mg", "1.5mg", "3mg", "4.5mg"],
  "frequency": "weekly",
  "cv_trial": "REWIND",
  "device": "single_use_pen"
},
"liraglutide": {
  "doses": ["0.6mg", "1.2mg", "1.8mg"],
  "frequency": "daily",
  "cv_trial": "LEADER",
  "cv_death_reduction": "22%"
},
"tirzepatide": {
  "doses": ["2.5mg", "5mg", "7.5mg", "10mg", "12.5mg", "15mg"],
  "frequency": "weekly",
  "mechanism": "dual GIP/GLP-1",
  "hba1c_reduction": "up to 2.5%",
  "weight_loss": "up to 22.5%"
}
}

```

## Clinical Decision Rules

### Insulin Initiation and Titration

```

class InsulinAlgorithms:
    basal_initiation = {
        "starting_dose": "10 units or 0.1-0.2 units/kg",
        "timing": "bedtime or same time daily",
        "titration": {
            "frequency": "every 3 days",
            "fasting_target": "80-130 mg/dL",
            "adjustment_rules": [
                {"if_fbg": ">180", "increase": 4},
                {"if_fbg": "140-180", "increase": 2},
                {"if_fbg": "80-130", "no_change": True},
                {"if_fbg": "<80", "decrease": 2}
            ]
        }
    }

```

```

    ]
  }
}

prandial_addition = {
  "when": "basal optimized but A1c above target",
  "starting": "4 units or 10% of basal dose",
  "meal": "largest meal first",
  "titration": "increase by 1-2 units every 3 days until target"
}

basal_bolus_optimization = {
  "basal_testing": "skip meal and check BG q2h",
  "dawn_phenomenon": "3am BG check",
  "somogyi_effect": "2-3am hypoglycemia followed by rebound"
}

```

## Cardiovascular Risk Management

### ASCVD Risk Calculation and Management

```

class CVRiskManagement:
    risk_calculator = {
        "variables": [
            "age", "sex", "race", "total_cholesterol",
            "hdl", "systolic_bp", "bp_treatment",
            "diabetes", "smoking"
        ],
        "categories": {
            "low": "<5%",
            "borderline": "5-7.4%",
            "intermediate": "7.5-19.9%",
            "high": "≥20%"
        }
    }

    statin_therapy = {
        "diabetes_40_75_years": "moderate_intensity",
        "diabetes_with_ascvd": "high_intensity",
        "diabetes_10yr_risk_>20%": "high_intensity",
        "high_intensity": ["atorvastatin 40-80mg", "rosuvastatin 20-40mg"],
        "moderate_intensity": ["atorvastatin 10-20mg", "rosuvastatin 5-10mg",
                                "simvastatin 20-40mg", "pravastatin 40-80mg"]
    }

```



```

}

bp_targets = {
  "general_diabetes": "<140/90",
  "high_cv_risk": "<130/80 if tolerated",
  "albuminuria": "<130/80"
}

antiplatelet = {
  "secondary_prevention": "aspirin 75-162mg",
  "primary_prevention": "consider if 10yr risk >10%"
}

```

## Microvascular Complication Screening

### Comprehensive Screening Protocols

```

class ComplicationScreening:
  retinopathy = {
    "initial": "at diagnosis for T2D, 5 years after onset for T1D",
    "frequency": "annual if no retinopathy",
    "method": "dilated funduscopy or retinal photography",
    "referral": "moderate NPDR or worse"
  }

  nephropathy = {
    "tests": ["annual eGFR", "annual UACR"],
    "ckd_staging": {
      "G1": {"egfr": "≥90", "description": "normal"},
      "G2": {"egfr": "60-89", "description": "mild decrease"},
      "G3a": {"egfr": "45-59", "description": "moderate decrease"},
      "G3b": {"egfr": "30-44", "description": "moderate decrease"},
      "G4": {"egfr": "15-29", "description": "severe decrease"},
      "G5": {"egfr": "<15", "description": "kidney failure"}
    },
    "albuminuria_categories": {
      "A1": "<30 mg/g",
      "A2": "30-300 mg/g",
      "A3": ">300 mg/g"
    }
  }

  neuropathy = {

```

```
"screening": "annually starting at diagnosis for T2D",
"tests": [
  "10-g monofilament",
  "vibration perception",
  "ankle reflexes",
  "temperature sensation"
],
"autonomic": [
  "orthostatic hypotension",
  "gastroparesis symptoms",
  "erectile dysfunction"
]
}
```

---

# Document 4: Technical Implementation Guide

## 4.1 Development Environment Setup

### Prerequisites and Tools

#### Development Machine Requirements

##### # Minimum hardware requirements

- CPU: 8 cores (16 threads)
- RAM: 32GB (64GB recommended)
- Storage: 500GB SSD
- GPU: NVIDIA RTX 3060 or better (for local ML development)

##### # Required software

- Docker Desktop 4.26+
- Kubernetes (minikube or kind for local)
- Python 3.11+
- Node.js 20 LTS
- Go 1.21+
- Rust 1.75+
- PostgreSQL 16
- Redis 7+

## Local Development Setup

### 1. Clone and Initialize Repository

```
# Clone repository
git clone https://github.com/easy-diabetes/easy-diabetes.git
cd easy-diabetes

# Initialize git-lfs for large model files
git lfs install
git lfs pull

# Setup pre-commit hooks
pip install pre-commit
pre-commit install

# Copy environment templates
cp .env.example .env.local
cp docker-compose.override.yml.example docker-compose.override.yml
```

### 2. Container Environment Setup

```
# docker-compose.yml
version: '3.9'

services:
  postgres:
    image: timescale/timescaledb-ha:pg16
    environment:
      POSTGRES_USER: easy_diabetes
      POSTGRES_PASSWORD: ${DB_PASSWORD}
      POSTGRES_DB: easy_diabetes
    volumes:
      - postgres_data:/var/lib/postgresql/data
      - ./init-scripts:/docker-entrypoint-initdb.d
    ports:
      - "5432:5432"

  neo4j:
    image: neo4j:5.15-enterprise
    environment:
      NEO4J_AUTH: neo4j/${NEO4J_PASSWORD}
      NEO4J_ACCEPT_LICENSE_AGREEMENT: yes
      NEO4J_dbms_memory_heap_max__size: 2G
```

volumes:  
- neo4j\_data:/data

ports:  
- "7474:7474"  
- "7687:7687"

redis:  
image: redis:7-alpine  
command: redis-server --appendonly yes  
volumes:  
- redis\_data:/data  
ports:  
- "6379:6379"

kafka:  
image: confluentinc/cp-kafka:7.5.0  
depends\_on:  
- zookeeper  
environment:  
KAFKA\_BROKER\_ID: 1  
KAFKA\_ZOOKEEPER\_CONNECT: zookeeper:2181  
KAFKA\_ADVERTISED\_LISTENERS: PLAINTEXT://localhost:9092  
KAFKA\_OFFSETS\_TOPIC\_REPLICATION\_FACTOR: 1  
ports:  
- "9092:9092"

clinical-api:  
build:  
context: ./services/clinical-api  
dockerfile: Dockerfile.dev  
environment:  
DATABASE\_URL:  
postgresql://easy\_diabetes:\${DB\_PASSWORD}@postgres:5432/easy\_diabetes  
REDIS\_URL: redis://redis:6379  
NEO4J\_URI: bolt://neo4j:7687  
volumes:  
- ./services/clinical-api:/app  
ports:  
- "8000:8000"  
command: uvicorn main:app --reload --host 0.0.0.0

volumes:  
postgres\_data:  
neo4j\_data:

redis\_data:

### 3. Python Environment Setup

# Create virtual environment

```
python -m venv venv
```

```
source venv/bin/activate # On Windows: venv\Scripts\activate
```

# Install development dependencies

```
pip install -r requirements-dev.txt
```

# Install project in editable mode

```
pip install -e .
```

# Setup Jupyter kernel for notebooks

```
python -m ipykernel install --user --name easy-diabetes
```

# Install ML dependencies

```
pip install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118
```

```
pip install transformers datasets accelerate
```

### 4. Frontend Development Setup

# Navigate to frontend directory

```
cd frontend
```

# Install dependencies

```
npm install
```

# Setup environment variables

```
cp .env.example .env.local
```

# Start development server

```
npm run dev
```

# In another terminal, start Storybook for component development

```
npm run storybook
```

## IDE Configuration

### VS Code Settings

```
{
  "python.linting.enabled": true,
  "python.linting.pylintEnabled": true,
  "python.linting.flake8Enabled": true,
  "python.formatting.provider": "black",
  "python.formatting.blackArgs": ["--line-length", "100"],
  "python.sortImports.provider": "isort",
  "python.testing.pytestEnabled": true,
  "editor.formatOnSave": true,
  "editor.codeActionsOnSave": {
    "source.organizeImports": true
  },
  "[typescript]": {
    "editor.defaultFormatter": "esbenp.prettier-vscode"
  },
  "[typescriptreact]": {
    "editor.defaultFormatter": "esbenp.prettier-vscode"
  }
}
```

## 4.2 Core Service Implementation

### Clinical Decision Service

#### FastAPI Application Structure

```
# services/clinical-api/main.py
from fastapi import FastAPI, Depends, HTTPException
from fastapi.middleware.cors import CORSMiddleware
from contextlib import asynccontextmanager
import uvicorn

from app.core.config import settings
from app.core.database import init_db
from app.api.v1.api import api_router
from app.middleware.auth import AuthMiddleware
from app.middleware.logging import LoggingMiddleware
from app.core.events import create_start_app_handler, create_stop_app_handler

@asynccontextmanager
async def lifespan(app: FastAPI):
    # Startup
    await create_start_app_handler(app)()
```

```

    await init_db()
    yield
    # Shutdown
    await create_stop_app_handler(app)()

app = FastAPI(
    title="EASY-Diabetes Clinical API",
    description="Clinical Decision Support System for Diabetes Management",
    version="2.0.0",
    lifespan=lifespan
)

# Middleware
app.add_middleware(
    CORSMiddleware,
    allow_origins=settings.BACKEND_CORS_ORIGINS,
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)
app.add_middleware(AuthMiddleware)
app.add_middleware(LoggingMiddleware)

# Include routers
app.include_router(api_router, prefix=settings.API_V1_STR)

if __name__ == "__main__":
    uvicorn.run(
        "main:app",
        host="0.0.0.0",
        port=8000,
        reload=settings.DEBUG,
        log_config="logging.yaml"
    )

```

## Clinical Decision Engine Implementation

```

# services/clinical-api/app/core/clinical_engine.py
from typing import List, Dict, Any, Optional
import asyncio
from datetime import datetime, timedelta

from app.models.patient import Patient
from app.models.clinical import ClinicalData, Medication, Condition

```

```

from app.services.guidelines import GuidelineService
from app.services.risk_assessment import RiskAssessmentService
from app.services.medication_optimizer import MedicationOptimizer
from app.ai.agents import ClinicalAgentOrchestrator

class ClinicalDecisionEngine:
    def __init__(self):
        self.guideline_service = GuidelineService()
        self.risk_service = RiskAssessmentService()
        self.medication_optimizer = MedicationOptimizer()
        self.agent_orchestrator = ClinicalAgentOrchestrator()

    async def generate_recommendations(
        self,
        patient: Patient,
        clinical_data: ClinicalData,
        request_context: Dict[str, Any]
    ) -> Dict[str, Any]:
        """Generate comprehensive clinical recommendations"""

        # Parallel execution of assessment tasks
        assessment_tasks = [
            self._assess_current_state(patient, clinical_data),
            self._calculate_risk_scores(patient, clinical_data),
            self._identify_treatment_gaps(patient, clinical_data),
            self._check_medication_issues(patient, clinical_data)
        ]

        results = await asyncio.gather(*assessment_tasks)
        current_state, risk_scores, treatment_gaps, medication_issues = results

        # Generate recommendations using AI agents
        agent_recommendations = await self.agent_orchestrator.generate_recommendations(
            patient=patient,
            clinical_data=clinical_data,
            assessment_results={
                "current_state": current_state,
                "risk_scores": risk_scores,
                "treatment_gaps": treatment_gaps,
                "medication_issues": medication_issues
            }
        )

        # Validate recommendations against guidelines

```



```

validated_recommendations = await self._validate_recommendations(
    agent_recommendations,
    patient,
    clinical_data
)

# Generate explanation and evidence
explanation = await self._generate_explanation(
    validated_recommendations,
    patient,
    clinical_data
)

return {
    "recommendations": validated_recommendations,
    "risk_assessment": risk_scores,
    "explanation": explanation,
    "evidence": self._compile_evidence(validated_recommendations),
    "confidence_scores": self._calculate_confidence(validated_recommendations),
    "generated_at": datetime.utcnow().isoformat()
}

async def _assess_current_state(
    self,
    patient: Patient,
    clinical_data: ClinicalData
) -> Dict[str, Any]:
    """Assess patient's current clinical state"""

    latest_labs = clinical_data.get_recent_labs(days=90)
    current_medications = clinical_data.get_active_medications()
    recent_events = clinical_data.get_recent_events(days=30)

    return {
        "glycemic_control": self._assess_glycemic_control(latest_labs),
        "medication_adherence": await self._calculate_adherence(patient, current_medications),
        "complication_status": self._assess_complications(clinical_data),
        "lifestyle_factors": await self._assess_lifestyle(patient),
        "recent_events": self._summarize_events(recent_events)
    }

```

## AI Agent Implementation

## Multi-Agent Orchestrator

```
# services/clinical-api/app/ai/agents/orchestrator.py
from typing import Dict, Any, List
import asyncio
from langchain.agents import AgentExecutor
from langchain.chains import LLMChain
from langchain.memory import ConversationSummaryBufferMemory

from app.ai.agents.clinical_assessment import ClinicalAssessmentAgent
from app.ai.agents.guideline_compliance import GuidelineComplianceAgent
from app.ai.agents.medication_optimization import MedicationOptimizationAgent
from app.ai.agents.safety_validation import SafetyValidationAgent
from app.ai.agents.explanation_generation import ExplanationGenerationAgent

class ClinicalAgentOrchestrator:
    def __init__(self):
        self.agents = {
            "assessment": ClinicalAssessmentAgent(),
            "guidelines": GuidelineComplianceAgent(),
            "medication": MedicationOptimizationAgent(),
            "safety": SafetyValidationAgent(),
            "explanation": ExplanationGenerationAgent()
        }

        self.memory = ConversationSummaryBufferMemory(
            memory_key="chat_history",
            return_messages=True,
            max_token_limit=2000
        )

    async def generate_recommendations(
        self,
        patient: Dict[str, Any],
        clinical_data: Dict[str, Any],
        assessment_results: Dict[str, Any]
    ) -> Dict[str, Any]:
        """Orchestrate multiple agents to generate recommendations"""

        # Phase 1: Clinical Assessment
        assessment_output = await self.agents["assessment"].analyze(
            patient=patient,
            clinical_data=clinical_data,
            context=assessment_results
```

```

)

# Phase 2: Guideline Application (parallel execution)
guideline_tasks = []
for condition in patient.get("conditions", []):
    task = self.agents["guidelines"].apply_guidelines(
        condition=condition,
        patient_context=patient,
        assessment=assessment_output
    )
    guideline_tasks.append(task)

guideline_results = await asyncio.gather(*guideline_tasks)

# Phase 3: Medication Optimization
medication_recommendations = await self.agents["medication"].optimize(
    current_medications=clinical_data.get("medications", []),
    patient_context=patient,
    guidelines=guideline_results,
    assessment=assessment_output
)

# Phase 4: Safety Validation
safety_validated = await self.agents["safety"].validate(
    recommendations=medication_recommendations,
    patient=patient,
    clinical_data=clinical_data
)

# Phase 5: Generate Explanations
explanations = await self.agents["explanation"].generate(
    recommendations=safety_validated,
    reasoning_chain=self._build_reasoning_chain(
        assessment_output,
        guideline_results,
        medication_recommendations,
        safety_validated
    )
)

return {
    "recommendations": safety_validated["safe_recommendations"],
    "explanations": explanations,
    "confidence": self._calculate_overall_confidence(safety_validated),
}

```

```

    "agent_outputs": {
        "assessment": assessment_output,
        "guidelines": guideline_results,
        "medication": medication_recommendations,
        "safety": safety_validated
    }
}

```

## Clinical Assessment Agent

```

# services/clinical-api/app/ai/agents/clinical_assessment.py
from langchain.agents import Tool, AgentExecutor
from langchain.chains import LLMChain
from langchain.prompts import PromptTemplate
from langchain.llms import ChatOpenAI

class ClinicalAssessmentAgent:
    def __init__(self):
        self.llm = ChatOpenAI(
            model="gpt-4-turbo",
            temperature=0.1,
            max_tokens=2000
        )

        self.tools = [
            Tool(
                name="LabAnalyzer",
                func=self._analyze_labs,
                description="Analyze laboratory results and identify abnormalities"
            ),
            Tool(
                name="CGMAAnalyzer",
                func=self._analyze_cgm,
                description="Analyze continuous glucose monitoring data patterns"
            ),
            Tool(
                name="MedicationReviewer",
                func=self._review_medications,
                description="Review current medications for efficacy and issues"
            ),
            Tool(
                name="ComplicationScreener",
                func=self._screen_complications,
                description="Screen for diabetes complications"
            )

```

```
)  
]
```

```
self.assessment_prompt = PromptTemplate(  
    input_variables=["patient_data", "clinical_data", "tool_outputs"],  
    template="""  
    You are an expert endocrinologist assessing a patient with diabetes.
```

```
  
    Patient Information:  
    {patient_data}
```

```
  
    Clinical Data:  
    {clinical_data}
```

```
  
    Analysis Results:  
    {tool_outputs}
```

```
  
    Provide a comprehensive clinical assessment including:  
    1. Current glycemic control status  
    2. Risk factors and complications  
    3. Treatment effectiveness  
    4. Key clinical concerns  
    5. Recommended focus areas for optimization
```

```
  
    Format your response as a structured clinical assessment.  
    """
```

```
)
```

```
async def analyze(  
    self,  
    patient: Dict[str, Any],  
    clinical_data: Dict[str, Any],  
    context: Dict[str, Any]  
) -> Dict[str, Any]:  
    """Perform comprehensive clinical assessment"""
```

```
  
    # Execute tools to gather information  
    tool_results = {}  
    for tool in self.tools:  
        result = await tool.func(clinical_data)  
        tool_results[tool.name] = result
```

```
  
    # Generate assessment using LLM  
    assessment_chain = LLMChain(  

```

```

        llm=self.llm,
        prompt=self.assessment_prompt
    )

    assessment = await assessment_chain.arun(
        patient_data=patient,
        clinical_data=clinical_data,
        tool_outputs=tool_results
    )

    # Structure the output
    return {
        "clinical_summary": assessment,
        "tool_results": tool_results,
        "key_findings": self._extract_key_findings(assessment),
        "risk_factors": self._identify_risk_factors(tool_results),
        "optimization_opportunities": self._identify_opportunities(assessment)
    }

```

## Database Schema Implementation

### PostgreSQL Schema

```

-- Core patient schema
CREATE SCHEMA IF NOT EXISTS patient;
CREATE SCHEMA IF NOT EXISTS clinical;
CREATE SCHEMA IF NOT EXISTS audit;

-- Enable necessary extensions
CREATE EXTENSION IF NOT EXISTS "uuid-ossf";
CREATE EXTENSION IF NOT EXISTS "pgcrypto";
CREATE EXTENSION IF NOT EXISTS "btree_gin";

-- Patient table with encryption for sensitive data
CREATE TABLE patient.patients (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    external_id VARCHAR(100) UNIQUE NOT NULL,
    -- Encrypted personal information
    first_name TEXT NOT NULL,
    last_name TEXT NOT NULL,
    date_of_birth DATE NOT NULL,
    gender VARCHAR(10) CHECK (gender IN ('male', 'female', 'other')),
    -- Contact information (encrypted)

```

```

email TEXT,
phone TEXT,
-- Clinical identifiers
mrn VARCHAR(50),
-- Metadata
created_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP,
created_by UUID REFERENCES auth.users(id),
is_active BOOLEAN DEFAULT true,
-- Indexes
CONSTRAINT patients_external_id_key UNIQUE (external_id)
);

-- Enable row-level security
ALTER TABLE patient.patients ENABLE ROW LEVEL SECURITY;

-- Clinical measurements with partitioning
CREATE TABLE clinical.measurements (
  id UUID DEFAULT uuid_generate_v4(),
  patient_id UUID NOT NULL REFERENCES patient.patients(id),
  measurement_type VARCHAR(50) NOT NULL,
  value NUMERIC(10,3) NOT NULL,
  unit VARCHAR(20) NOT NULL,
  measured_at TIMESTAMPTZ NOT NULL,
  source VARCHAR(50), -- 'manual', 'device', 'lab', 'ehr'
  device_id UUID,
  metadata JSONB DEFAULT '{}',
  created_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP,
  PRIMARY KEY (id, measured_at)
) PARTITION BY RANGE (measured_at);

-- Create monthly partitions
DO $$
DECLARE
  start_date DATE := '2024-01-01';
  end_date DATE;
  partition_name TEXT;
BEGIN
  FOR i IN 0..23 LOOP
    end_date := start_date + INTERVAL '1 month';
    partition_name := 'measurements_' || TO_CHAR(start_date, 'YYYY_MM');

    EXECUTE format(
      'CREATE TABLE clinical.%I PARTITION OF clinical.measurements

```

```

        FOR VALUES FROM (%L) TO (%L)',
        partition_name, start_date, end_date
    );

    start_date := end_date;
END LOOP;
END $$;

-- Medications table
CREATE TABLE clinical.medications (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    patient_id UUID NOT NULL REFERENCES patient.patients(id),
    medication_name VARCHAR(200) NOT NULL,
    generic_name VARCHAR(200),
    drug_class VARCHAR(100),
    dose NUMERIC(10,3),
    dose_unit VARCHAR(20),
    frequency VARCHAR(50),
    route VARCHAR(20),
    start_date DATE NOT NULL,
    end_date DATE,
    discontinued_reason TEXT,
    prescriber_id UUID,
    is_active BOOLEAN DEFAULT true,
    created_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP
);

-- Create indexes
CREATE INDEX idx_medications_patient_active ON clinical.medications(patient_id, is_active);
CREATE INDEX idx_medications_drug_class ON clinical.medications(drug_class);

-- Conditions table
CREATE TABLE clinical.conditions (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    patient_id UUID NOT NULL REFERENCES patient.patients(id),
    condition_code VARCHAR(20),
    condition_system VARCHAR(20), -- 'ICD10', 'SNOMED'
    condition_display TEXT NOT NULL,
    clinical_status VARCHAR(20), -- 'active', 'resolved', 'inactive'
    verification_status VARCHAR(20), -- 'confirmed', 'provisional', 'differential'
    onset_date DATE,
    abatement_date DATE,
    recorded_date DATE NOT NULL,

```



```

    recorder_id UUID,
    notes TEXT,
    created_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP
);

-- Laboratory results
CREATE TABLE clinical.lab_results (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    patient_id UUID NOT NULL REFERENCES patient.patients(id),
    lab_type VARCHAR(50) NOT NULL,
    loinc_code VARCHAR(20),
    value VARCHAR(50) NOT NULL,
    value_numeric NUMERIC(10,3),
    unit VARCHAR(20),
    reference_range VARCHAR(50),
    interpretation VARCHAR(20), -- 'normal', 'high', 'low', 'critical'
    collected_at TIMESTAMPTZ NOT NULL,
    resulted_at TIMESTAMPTZ,
    performing_lab VARCHAR(100),
    created_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP
);

-- Create indexes for common queries
CREATE INDEX idx_lab_results_patient_type_date ON clinical.lab_results(patient_id, lab_type,
collected_at DESC);
CREATE INDEX idx_lab_results_interpretation ON clinical.lab_results(interpretation) WHERE
interpretation != 'normal';

-- Audit log table
CREATE TABLE audit.audit_log (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    user_id UUID NOT NULL,
    patient_id UUID,
    action VARCHAR(50) NOT NULL,
    resource_type VARCHAR(50) NOT NULL,
    resource_id UUID,
    changes JSONB,
    ip_address INET,
    user_agent TEXT,
    created_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP
);

-- Create indexes for audit queries
CREATE INDEX idx_audit_log_user_date ON audit.audit_log(user_id, created_at DESC);

```

```
CREATE INDEX idx_audit_log_patient_date ON audit.audit_log(patient_id, created_at DESC);
```

## Neo4j Knowledge Graph Schema

```
// Create constraints for unique identifiers
```

```
CREATE CONSTRAINT patient_id IF NOT EXISTS ON (p:Patient) ASSERT p.id IS UNIQUE;  
CREATE CONSTRAINT medication_id IF NOT EXISTS ON (m:Medication) ASSERT m.id IS  
UNIQUE;
```

```
CREATE CONSTRAINT condition_id IF NOT EXISTS ON (c:Condition) ASSERT c.id IS  
UNIQUE;
```

```
CREATE CONSTRAINT guideline_id IF NOT EXISTS ON (g:Guideline) ASSERT g.id IS  
UNIQUE;
```

```
// Create indexes for performance
```

```
CREATE INDEX patient_external_id IF NOT EXISTS FOR (p:Patient) ON (p.external_id);
```

```
CREATE INDEX medication_name IF NOT EXISTS FOR (m:Medication) ON (m.name);
```

```
CREATE INDEX condition_code IF NOT EXISTS FOR (c:Condition) ON (c.code);
```

```
// Example knowledge graph structure
```

```
// Create patient node
```

```
CREATE (p:Patient {  
  id: 'uuid-here',  
  external_id: 'EXT123',  
  age_group: '45-54',  
  diabetes_type: 'type2',  
  duration_years: 5  
})
```

```
// Create condition nodes
```

```
CREATE (dm:Condition {  
  id: 'uuid-dm',  
  code: 'E11.9',  
  system: 'ICD10',  
  name: 'Type 2 diabetes mellitus without complications'  
})
```

```
CREATE (htn:Condition {  
  id: 'uuid-htn',  
  code: 'I10',  
  system: 'ICD10',  
  name: 'Essential hypertension'  
})
```

```
// Create relationships
```

```
CREATE (p)-[:HAS_CONDITION {diagnosed_date: date('2019-01-15')}]>(dm)
CREATE (p)-[:HAS_CONDITION {diagnosed_date: date('2020-03-20')}]>(htn)
```

```
// Create medication nodes and relationships
CREATE (met:Medication {
  id: 'uuid-met',
  name: 'Metformin',
  drug_class: 'Biguanide',
  mechanism: 'Decreases hepatic glucose production'
})
```

```
CREATE (p)-[:TAKES_MEDICATION {
  dose: 1000,
  unit: 'mg',
  frequency: 'BID',
  start_date: date('2019-02-01'),
  adherence_rate: 0.85
}]>(met)
```

```
// Create similar patient relationships for collaborative filtering
MATCH (p1:Patient), (p2:Patient)
WHERE p1.id <> p2.id
  AND abs(p1.age_group - p2.age_group) <= 1
  AND p1.diabetes_type = p2.diabetes_type
  AND abs(p1.duration_years - p2.duration_years) <= 2
CREATE (p1)-[:SIMILAR_TO {
  similarity_score: 0.85,
  calculated_at: datetime()
}]>(p2)
```

```
// Guideline relationships
CREATE (g:Guideline {
  id: 'uuid-guide',
  name: 'ADA Standards of Care 2024',
  version: '2024.1',
  section: 'Pharmacologic Approaches'
})
```

```
CREATE (met)-[:RECOMMENDED_BY {
  recommendation_strength: 'strong',
  evidence_level: 'A'
}]>(g)
```

```
// Contraindication relationships
```

```

CREATE (ckd:Condition {
  id: 'uuid-ckd',
  code: 'N18.5',
  name: 'Chronic kidney disease, stage 5'
})

CREATE (met)-[:CONTRAINDICATED_IN {
  absolute: true,
  reason: 'eGFR < 30'
}]->(ckd)

```

## API Endpoint Implementation

### Patient Management Endpoints

```

# services/clinical-api/app/api/v1/endpoints/patients.py
from typing import List, Optional
from fastapi import APIRouter, Depends, HTTPException, Query
from sqlalchemy.ext.asyncio import AsyncSession

from app.api.deps import get_db, get_current_user
from app.models.user import User
from app.schemas.patient import (
    PatientCreate,
    PatientUpdate,
    PatientResponse,
    PatientListResponse
)
from app.crud.patient import patient_crud
from app.core.security import check_permission

router = APIRouter()

@router.post("/", response_model=PatientResponse)
async def create_patient(
    *,
    db: AsyncSession = Depends(get_db),
    patient_in: PatientCreate,
    current_user: User = Depends(get_current_user)
) -> PatientResponse:
    """Create new patient record"""

    # Check permissions

```

```

if not check_permission(current_user, "patient:create"):
    raise HTTPException(status_code=403, detail="Insufficient permissions")

# Check for duplicate
existing = await patient_crud.get_by_external_id(
    db, external_id=patient_in.external_id
)
if existing:
    raise HTTPException(
        status_code=400,
        detail="Patient with this external ID already exists"
    )

# Create patient
patient = await patient_crud.create_with_owner(
    db, obj_in=patient_in, owner_id=current_user.id
)

# Log audit event
await audit_log.create_entry(
    db,
    user_id=current_user.id,
    action="create_patient",
    resource_type="patient",
    resource_id=patient.id
)

return patient

@router.get("/{patient_id}", response_model=PatientResponse)
async def get_patient(
    *,
    db: AsyncSession = Depends(get_db),
    patient_id: str,
    current_user: User = Depends(get_current_user)
) -> PatientResponse:
    """Get patient by ID"""

    # Get patient
    patient = await patient_crud.get(db, id=patient_id)
    if not patient:
        raise HTTPException(status_code=404, detail="Patient not found")

    # Check access permissions

```

```

if not check_patient_access(current_user, patient):
    raise HTTPException(status_code=403, detail="Access denied")

return patient

@router.get("/", response_model=PatientListResponse)
async def list_patients(
    db: AsyncSession = Depends(get_db),
    skip: int = Query(0, ge=0),
    limit: int = Query(100, ge=1, le=1000),
    search: Optional[str] = None,
    current_user: User = Depends(get_current_user)
) -> PatientListResponse:
    """List patients with pagination and search"""

    # Check permissions
    if not check_permission(current_user, "patient:list"):
        raise HTTPException(status_code=403, detail="Insufficient permissions")

    # Get patients based on user role
    if current_user.role == "clinician":
        # Only show assigned patients
        patients = await patient_crud.get_by_clinician(
            db,
            clinician_id=current_user.id,
            skip=skip,
            limit=limit,
            search=search
        )
    else:
        # Admin can see all
        patients = await patient_crud.get_multi(
            db, skip=skip, limit=limit, search=search
        )

    return PatientListResponse(
        items=patients,
        total=len(patients),
        skip=skip,
        limit=limit
    )

```

## Clinical Recommendations Endpoint

```

# services/clinical-api/app/api/v1/endpoints/recommendations.py
from fastapi import APIRouter, Depends, HTTPException, BackgroundTasks
from sqlalchemy.ext.asyncio import AsyncSession
import asyncio

from app.api.deps import get_db, get_current_user
from app.models.user import User
from app.schemas.recommendations import (
    RecommendationRequest,
    RecommendationResponse,
    RecommendationFeedback
)
from app.services.clinical_engine import ClinicalDecisionEngine
from app.services.cache import recommendation_cache
from app.core.monitoring import metrics

router = APIRouter()

@router.post("/{patient_id}/generate", response_model=RecommendationResponse)
async def generate_recommendations(
    *,
    db: AsyncSession = Depends(get_db),
    patient_id: str,
    request: RecommendationRequest,
    background_tasks: BackgroundTasks,
    current_user: User = Depends(get_current_user)
) -> RecommendationResponse:
    """Generate clinical recommendations for patient"""

    # Start timing for metrics
    start_time = asyncio.get_event_loop().time()

    try:
        # Check patient access
        patient = await get_patient_with_access_check(
            db, patient_id, current_user
        )

        # Check cache first
        cache_key = f"recommendations:{patient_id}:{request.hash()}"
        cached = await recommendation_cache.get(cache_key)
        if cached and not request.force_refresh:
            metrics.cache_hits.inc()
            return cached

```

```

# Get clinical data
clinical_data = await get_clinical_data(db, patient_id)

# Initialize decision engine
engine = ClinicalDecisionEngine()

# Generate recommendations
recommendations = await engine.generate_recommendations(
    patient=patient,
    clinical_data=clinical_data,
    request_context={
        "user_id": current_user.id,
        "request_params": request.dict(),
        "timestamp": datetime.utcnow()
    }
)

# Cache results
await recommendation_cache.set(
    cache_key,
    recommendations,
    expire=300 # 5 minutes
)

# Background tasks
background_tasks.add_task(
    log_recommendation_generation,
    patient_id=patient_id,
    user_id=current_user.id,
    recommendations=recommendations
)

# Record metrics
duration = asyncio.get_event_loop().time() - start_time
metrics.recommendation_generation_time.observe(duration)
metrics.recommendations_generated.inc()

return RecommendationResponse(**recommendations)

except Exception as e:
    metrics.recommendation_errors.inc()
    raise HTTPException(
        status_code=500,

```



```

        detail=f"Error generating recommendations: {str(e)}"
    )

@router.post("/{patient_id}/recommendations/{recommendation_id}/feedback")
async def submit_recommendation_feedback(
    *,
    db: AsyncSession = Depends(get_db),
    patient_id: str,
    recommendation_id: str,
    feedback: RecommendationFeedback,
    current_user: User = Depends(get_current_user)
) -> dict:
    """Submit feedback on a recommendation"""

    # Verify access
    patient = await get_patient_with_access_check(
        db, patient_id, current_user
    )

    # Store feedback
    await store_recommendation_feedback(
        db,
        patient_id=patient_id,
        recommendation_id=recommendation_id,
        feedback=feedback,
        user_id=current_user.id
    )

    # Update ML models if accepted/rejected
    if feedback.action in ["accepted", "rejected"]:
        await queue_model_update(
            recommendation_id=recommendation_id,
            action=feedback.action,
            reason=feedback.reason
        )

    return {"status": "feedback_recorded"}

```

## Frontend Implementation

### React Component Architecture

```
// frontend/src/components/clinical/ClinicalDashboard.tsx
```

```

import React, { useState, useEffect } from 'react'
import { useQuery, useMutation } from '@tanstack/react-query'
import {
  Card,
  CardContent,
  CardDescription,
  CardHeader,
  CardTitle
} from '@components/ui/card'
import { Button } from '@components/ui/button'
import { Alert, AlertDescription } from '@components/ui/alert'
import { Tabs, TabsContent, TabsList, TabsTrigger } from '@components/ui/tabs'

import { PatientOverview } from './PatientOverview'
import { GlycemicControl } from './GlycemicControl'
import { MedicationManagement } from './MedicationManagement'
import { RiskAssessment } from './RiskAssessment'
import { ClinicalRecommendations } from './ClinicalRecommendations'
import { usePatientStore } from '@stores/patientStore'
import { api } from '@lib/api'

interface ClinicalDashboardProps {
  patientId: string
}

export const ClinicalDashboard: React.FC<ClinicalDashboardProps> = ({
  patientId
}) => {
  const [activeTab, setActiveTab] = useState('overview')
  const { setCurrentPatient } = usePatientStore()

  // Fetch patient data
  const { data: patient, isLoading: patientLoading } = useQuery({
    queryKey: ['patient', patientId],
    queryFn: () => api.patients.get(patientId),
    onSuccess: (data) => setCurrentPatient(data)
  })

  // Fetch clinical data
  const { data: clinicalData, isLoading: clinicalLoading } = useQuery({
    queryKey: ['clinical-data', patientId],
    queryFn: () => api.clinical.getData(patientId),
    refetchInterval: 5 * 60 * 1000 // Refetch every 5 minutes
  })

```

```

// Generate recommendations mutation
const generateRecommendations = useMutation({
  mutationFn: () => api.recommendations.generate(patientId),
  onSuccess: () => {
    queryClient.invalidateQueries(['recommendations', patientId])
  }
})

if (patientLoading || clinicalLoading) {
  return <DashboardSkeleton />
}

return (
  <div className="space-y-6">
    {/* Patient Header */}
    <Card>
      <CardHeader>
        <div className="flex items-center justify-between">
          <div>
            <CardTitle className="text-2xl">
              {patient.firstName} {patient.lastName}
            </CardTitle>
            <CardDescription>
              MRN: {patient.mrn} | {calculateAge(patient.dateOfBirth)} years old
            </CardDescription>
          </div>
          <Button
            onClick={() => generateRecommendations.mutate()}
            disabled={generateRecommendations.isLoading}
          >
            {generateRecommendations.isLoading
              ? 'Generating...'
              : 'Generate Recommendations'}
          </Button>
        </div>
      </CardHeader>
    </Card>

    {/* Main Dashboard Tabs */}
    <Tabs value={activeTab} onValueChange={setActiveTab}>
      <TabsList className="grid w-full grid-cols-5">
        <TabsTrigger value="overview">Overview</TabsTrigger>
        <TabsTrigger value="glycemic">Glycemic Control</TabsTrigger>
      </TabsList>
    </Tabs>
  </div>
)

```

```

    <TabsTrigger value="medications">Medications</TabsTrigger>
    <TabsTrigger value="risk">Risk Assessment</TabsTrigger>
    <TabsTrigger value="recommendations">Recommendations</TabsTrigger>
  </TabsList>

  <TabsContent value="overview" className="space-y-4">
    <PatientOverview
      patient={patient}
      clinicalData={clinicalData}
    />
  </TabsContent>

  <TabsContent value="glycemic" className="space-y-4">
    <GlycemicControl
      patientId={patientId}
      clinicalData={clinicalData}
    />
  </TabsContent>

  <TabsContent value="medications" className="space-y-4">
    <MedicationManagement
      patientId={patientId}
      medications={clinicalData.medications}
    />
  </TabsContent>

  <TabsContent value="risk" className="space-y-4">
    <RiskAssessment
      patientId={patientId}
      riskScores={clinicalData.riskScores}
    />
  </TabsContent>

  <TabsContent value="recommendations" className="space-y-4">
    <ClinicalRecommendations
      patientId={patientId}
    />
  </TabsContent>
</Tabs>
</div>
)
}

```

## Glycemic Control Visualization Component

```
// frontend/src/components/clinical/GlycemicControl.tsx
import React from 'react'
import {
  LineChart,
  Line,
  AreaChart,
  Area,
  XAxis,
  YAxis,
  CartesianGrid,
  Tooltip,
  ResponsiveContainer,
  ReferenceLine
} from 'recharts'
import { Card, CardContent, CardHeader, CardTitle } from '@components/ui/card'
import { Badge } from '@components/ui/badge'

interface GlycemicControlProps {
  patientId: string
  clinicalData: ClinicalData
}

export const GlycemicControl: React.FC<GlycemicControlProps> = ({
  patientId,
  clinicalData
}) => {
  const latestHbA1c = clinicalData.labs.hba1c[0]
  const cgmData = processCGMData(clinicalData.cgm)
  const timeInRange = calculateTimeInRange(cgmData)

  return (
    <div className="grid gap-4 md:grid-cols-2 lg:grid-cols-3">
      { /* HbA1c Card */ }
      <Card>
        <CardHeader className="flex flex-row items-center justify-between space-y-0 pb-2">
          <CardTitle className="text-sm font-medium">HbA1c</CardTitle>
          <Badge variant={getHbA1cVariant(latestHbA1c.value)}>
            {getHbA1cStatus(latestHbA1c.value)}
          </Badge>
        </CardHeader>
        <CardContent>
          <div className="text-2xl font-bold">{latestHbA1c.value}%</div>
```

```

<p className="text-xs text-muted-foreground">
  {formatDate(latestHbA1c.date)}
</p>
<div className="mt-4">
  <ResponsiveContainer width="100%" height={100}>
    <LineChart data={clinicalData.labs.hba1c}>
      <Line
        type="monotone"
        dataKey="value"
        stroke="#8884d8"
        strokeWidth={2}
        dot={{ r: 4 }}
      />
      <ReferenceLine y={7} stroke="#ef4444" strokeDasharray="3 3" />
      <Tooltip />
    </LineChart>
  </ResponsiveContainer>
</div>
</CardContent>
</Card>

{/* Time in Range Card */}
<Card>
  <CardHeader>
    <CardTitle className="text-sm font-medium">Time in Range</CardTitle>
  </CardHeader>
  <CardContent>
    <div className="space-y-2">
      <TimeRangeBar
        label="Above Range (>180)"
        value={timeInRange.above}
        color="text-orange-500"
        target="<25%"
      />
      <TimeRangeBar
        label="In Range (70-180)"
        value={timeInRange.inRange}
        color="text-green-500"
        target=">70%"
      />
      <TimeRangeBar
        label="Below Range (<70)"
        value={timeInRange.below}
        color="text-red-500"

```

```

        target="<4%"
      />
    </div>
  </CardContent>
</Card>

{/* CGM Trace */}
<Card className="col-span-full">
  <CardHeader>
    <CardTitle>Continuous Glucose Monitor - 24 Hour View</CardTitle>
  </CardHeader>
  <CardContent>
    <ResponsiveContainer width="100%" height={300}>
      <AreaChart data={cgmData}>
        <CartesianGrid strokeDasharray="3 3" />
        <XAxis
          dataKey="time"
          tickFormatter={(time) => formatTime(time)}
        />
        <YAxis domain={[40, 400]} />
        <Tooltip
          labelFormatter={(time) => formatDateTime(time)}
          formatter={(value) => [`${value} mg/dL`, 'Glucose']}
        />

        {/* Target range background */}
        <ReferenceLine y={180} stroke="#ef4444" strokeDasharray="3 3" />
        <ReferenceLine y={70} stroke="#ef4444" strokeDasharray="3 3" />

        {/* Glucose trace */}
        <Area
          type="monotone"
          dataKey="glucose"
          stroke="#8884d8"
          fill="#8884d8"
          fillOpacity={0.3}
        />

        {/* Meal markers */}
        {cgmData.meals?.map((meal, idx) => (
          <ReferenceLine
            key={idx}
            x={meal.time}
            stroke="#10b981"

```

```
        strokeWidth={2}
        label="Meal"
      />
    )))
  </AreaChart>
</ResponsiveContainer>
</CardContent>
</Card>
</div>
)
}
```

---

# Document 5: Product Roadmap

## Phase 1: Foundation

**Goal:** Establish core infrastructure and basic clinical functionality

### Technical Milestones

- ☒ Core microservices architecture
- ☒ PostgreSQL and Neo4j setup
- ☒ Basic authentication/authorization
- ☒ FHIR R4 integration framework
- ☒ Initial API development

### Clinical Features

- ☒ Patient registration and management
- ☒ Basic medication tracking
- ☒ Lab result integration
- ☒ Simple clinical alerts
- ☒ Basic reporting

### Regulatory

- ☒ HIPAA compliance framework
- ☒ Initial FDA pre-submission meeting
- ☒ Quality management system setup



## Phase 2: AI Integration

**Goal:** Implement core AI capabilities and clinical decision support

### AI/ML Features

- ☐ Multi-agent architecture implementation
- ☐ Clinical assessment agent
- ☐ Guideline compliance agent
- ☐ Basic recommendation generation
- ☐ Explanation framework

### Clinical Enhancements

- ☐ CGM integration (Dexcom, FreeStyle Libre)
- ☐ Insulin dosing calculator
- ☐ Hypoglycemia prediction (2-hour)
- ☐ Medication optimization suggestions

### Integration

- ☐ Epic FHIR integration
- ☐ Cerner PowerChart integration
- ☐ Apple HealthKit integration

## Phase 3: Advanced Analytics

**Goal:** Sophisticated predictive models and personalization

### Predictive Analytics

- ☐ 4-hour glucose prediction
- ☐ HbA1c trajectory modeling
- ☐ Complication risk scoring
- ☐ Treatment response prediction

### Personalization

- ☐ Patient similarity matching
- ☐ Personalized targets
- ☐ Lifestyle factor integration
- ☐ Behavioral nudges

## Clinical Features

- ☐ Automated insulin adjustment
- ☐ Carb counting assistance
- ☐ Exercise impact prediction
- ☐ Sick day management

## Phase 4: Scale & Optimize

**Goal:** Production readiness and optimization

### Scalability

- ☐ Kubernetes autoscaling
- ☐ Multi-region deployment
- ☐ Performance optimization
- ☐ Load testing (1M users)

### Advanced Features

- ☐ Federated learning implementation
- ☐ Real-time collaboration tools
- ☐ Voice interface
- ☐ AR medication administration

### Regulatory

- ☐ FDA 510(k) submission
- ☐ CE marking process
- ☐ Clinical trial initiation

## Phase 5: Expansion

**Goal:** Broader disease coverage and global deployment

### Disease Expansion

- ☐ Type 1 diabetes full support
- ☐ Gestational diabetes module
- ☐ Prediabetes prevention
- ☐ Pediatric diabetes

### Geographic Expansion

- ☐ EU deployment
- ☐ Localization (10 languages)
- ☐ Regional guideline integration
- ☐ Local EHR integrations

## Advanced AI

- ☐ Multimodal learning
- ☐ Causal inference engine
- ☐ Automated clinical trials
- ☐ Digital twin modeling

## Long-term Vision

### Platform Evolution

- Comprehensive metabolic disease platform
- Integration with digital therapeutics
- Closed-loop automated care
- Population health management

### Technology Advancement

- Quantum computing for drug discovery
- Advanced robotics integration
- Brain-computer interfaces
- Nano-sensor integration

### Global Impact

- 10M+ patients managed
- 50+ country deployment
- 100+ health system integrations
- Measurable population health improvement

---

# Document 6: Project Knowledge Base

## 1. Domain Knowledge

## Diabetes Pathophysiology

### Type 2 Diabetes:

- Insulin resistance + relative insulin deficiency
- Progressive beta cell dysfunction
- Hepatic glucose overproduction
- Incretin deficiency/resistance
- Accelerated lipolysis
- Glucagon hypersecretion
- Renal glucose reabsorption increase
- Brain insulin resistance

### Type 1 Diabetes:

- Autoimmune beta cell destruction
- Absolute insulin deficiency
- Presence of autoantibodies (GAD, IA-2, ZnT8)
- Risk of diabetic ketoacidosis
- Usually juvenile onset (but LADA exists)

## Key Clinical Concepts

### Glycemic Targets:

- HbA1c: Average glucose over 3 months
- Time in Range: % time glucose 70-180 mg/dL
- Glycemic Variability: Coefficient of variation <36%
- Hypoglycemia: <70 mg/dL (Level 1), <54 mg/dL (Level 2)

### Complications:

#### Microvascular:

- Retinopathy (leading cause of blindness)
- Nephropathy (leading cause of ESRD)
- Neuropathy (leading cause of amputation)

#### Macrovascular:

- 2-4x increased CV risk
- Stroke risk doubled
- Peripheral arterial disease

#### Acute:

- Hypoglycemia
- Hyperglycemic crises (DKA, HHS)

## Treatment Principles

Lifestyle:

- Medical nutrition therapy
- 150 min/week moderate exercise
- Weight loss 5-7% if overweight
- Smoking cessation

Medications:

- Metformin first-line unless contraindicated
- Early combination therapy if HbA1c >1.5% above target
- Consider CV/renal benefits (SGLT2i, GLP-1)
- Avoid hypoglycemia in elderly
- Cost considerations

Monitoring:

- SMBG if on insulin or sulfonylureas
- CGM for all Type 1, intensive insulin
- HbA1c every 3 months if not at goal
- Annual complication screening

## 2. Technical Knowledge

### System Architecture Patterns

Microservices Benefits:

- Independent deployment
- Technology diversity
- Fault isolation
- Scalability
- Team autonomy

Event-Driven Architecture:

- Loose coupling
- Real-time processing
- Event sourcing
- CQRS pattern
- Audit trail

API Design:

- RESTful principles
- GraphQL for complex queries
- gRPC for internal services
- OpenAPI documentation
- Versioning strategy

## **AI/ML Concepts**

Multi-Agent Systems:

- Specialized agents for subtasks
- Orchestration layer
- Inter-agent communication
- Consensus mechanisms
- Explanation aggregation

Clinical NLP:

- Named entity recognition (medications, conditions)
- Relation extraction
- Negation detection
- Temporal reasoning
- Section segmentation

Predictive Models:

- Time series forecasting
- Survival analysis
- Causal inference
- Uncertainty quantification
- Model interpretability

## **Security & Compliance**

HIPAA Requirements:

- Access controls
- Audit trails
- Encryption (at rest & transit)
- Minimum necessary
- Business Associate Agreements

FDA SaMD:

- Design controls
- Risk management (ISO 14971)
- Clinical validation
- Post-market surveillance
- Change control

Zero Trust Security:

- Never trust, always verify
- Least privilege access

- Micro-segmentation
- Continuous monitoring
- Identity-centric

## 3. Implementation Best Practices

### Code Quality

```
# Example: Type hints and documentation
from typing import List, Dict, Optional, Union
from datetime import datetime
from pydantic import BaseModel, Field
```

```
class GlucoseReading(BaseModel):
    """Represents a blood glucose measurement.
```

Attributes:

```
    value: Glucose value in mg/dL
    timestamp: When the reading was taken
    meal_tag: Optional meal context
    notes: Optional user notes
    """
```

```
    value: float = Field(..., ge=20, le=600, description="Glucose in mg/dL")
    timestamp: datetime
    meal_tag: Optional[str] = Field(None, regex="^(before|after)_(breakfast|lunch|dinner|snack)$")
    notes: Optional[str] = Field(None, max_length=500)
```

```
class Config:
    schema_extra = {
        "example": {
            "value": 120,
            "timestamp": "2024-01-15T08:30:00Z",
            "meal_tag": "before_breakfast",
            "notes": "Feeling good"
        }
    }
```

### Testing Strategies

```
# Example: Comprehensive test structure
import pytest
from unittest.mock import Mock, AsyncMock
from datetime import datetime, timedelta
```

```

class TestGlucosePrediction:
    """Test suite for glucose prediction functionality"""

    @pytest.fixture
    def mock_cgm_data(self):
        """Generate mock CGM data for testing"""
        base_time = datetime.now()
        return [
            {"timestamp": base_time - timedelta(minutes=5*i),
             "glucose": 120 + 10*np.sin(i/10)}
            for i in range(288) # 24 hours of data
        ]

    @pytest.mark.asyncio
    async def test_prediction_accuracy(self, mock_cgm_data):
        """Test that predictions are within acceptable range"""
        predictor = GlucosePredictor()
        predictions = await predictor.predict(mock_cgm_data, horizon=60)

        assert predictions is not None
        assert len(predictions) == 12 # 60 minutes / 5 minute intervals
        assert all(20 <= p <= 400 for p in predictions)

    @pytest.mark.parametrize("horizon,expected_length", [
        (30, 6),
        (60, 12),
        (120, 24),
        (240, 48)
    ])
    async def test_prediction_horizons(self, mock_cgm_data, horizon, expected_length):
        """Test predictions for different time horizons"""
        predictor = GlucosePredictor()
        predictions = await predictor.predict(mock_cgm_data, horizon=horizon)
        assert len(predictions) == expected_length

```

## Performance Optimization

```

# Example: Caching and async optimization
from functools import lru_cache
from typing import List
import asyncio
import aioredis

```



```

class OptimizedClinicalService:
    def __init__(self):
        self.redis = None
        self._cache = {}

    async def initialize(self):
        """Initialize Redis connection"""
        self.redis = await aioredis.create_redis_pool(
            'redis://localhost',
            encoding='utf-8'
        )

    @lru_cache(maxsize=1000)
    def calculate_risk_score(self, patient_id: str) -> float:
        """Calculate risk score with LRU caching"""
        # Expensive calculation cached in memory
        return self._complex_risk_calculation(patient_id)

    async def get_patient_data_batch(self, patient_ids: List[str]) -> List[dict]:
        """Fetch multiple patients efficiently"""
        # Create tasks for parallel execution
        tasks = [self.get_patient_data(pid) for pid in patient_ids]

        # Execute in parallel with semaphore to limit concurrency
        sem = asyncio.Semaphore(10)
        async def bounded_fetch(patient_id):
            async with sem:
                return await self.get_patient_data(patient_id)

        bounded_tasks = [bounded_fetch(pid) for pid in patient_ids]
        return await asyncio.gather(*bounded_tasks)

    async def get_patient_data(self, patient_id: str) -> dict:
        """Get patient data with Redis caching"""
        # Check Redis cache first
        cached = await self.redis.get(f"patient:{patient_id}")
        if cached:
            return json.loads(cached)

        # Fetch from database
        data = await self._fetch_from_db(patient_id)

        # Cache with expiration
        await self.redis.setex(

```

```

        f"patient:{patient_id}",
        300, # 5 minute TTL
        json.dumps(data)
    )

```

```

    return data

```

## Monitoring and Observability

```

# Example: Comprehensive monitoring setup
from prometheus_client import Counter, Histogram, Gauge
from opentelemetry import trace
from opentelemetry.exporter.jaeger import JaegerExporter
from opentelemetry.sdk.trace import TracerProvider
from opentelemetry.sdk.trace.export import BatchSpanProcessor

```

```

# Metrics
recommendation_counter = Counter(
    'clinical_recommendations_total',
    'Total number of clinical recommendations generated',
    ['recommendation_type', 'status']
)

```

```

recommendation_latency = Histogram(
    'clinical_recommendation_duration_seconds',
    'Time spent generating recommendations',
    ['recommendation_type']
)

```

```

active_users = Gauge(
    'active_users_total',
    'Number of active users in the system'
)

```

```

# Tracing
trace.set_tracer_provider(TracerProvider())
tracer = trace.get_tracer(__name__)

```

```

jaeger_exporter = JaegerExporter(
    agent_host_name="localhost",
    agent_port=6831,
)

```

```

span_processor = BatchSpanProcessor(jaeger_exporter)

```

```

trace.get_tracer_provider().add_span_processor(span_processor)

class MonitoredClinicalService:
    @recommendation_latency.time()
    async def generate_recommendation(self, patient_id: str) -> dict:
        """Generate recommendation with monitoring"""
        with tracer.start_as_current_span("generate_recommendation") as span:
            span.set_attribute("patient.id", patient_id)

            try:
                # Generate recommendation
                recommendation = await self._generate(patient_id)

                # Record metrics
                recommendation_counter.labels(
                    recommendation_type=recommendation['type'],
                    status='success'
                ).inc()

                span.set_attribute("recommendation.type", recommendation['type'])
                span.set_attribute("recommendation.confidence", recommendation['confidence'])

                return recommendation

            except Exception as e:
                recommendation_counter.labels(
                    recommendation_type='unknown',
                    status='error'
                ).inc()

                span.record_exception(e)
                span.set_status(trace.StatusCode.ERROR)
                raise

```

## 4. Troubleshooting Guide

### Common Issues and Solutions

#### Issue: High latency in recommendation generation

##### # Diagnosis steps

##### 1. Check service metrics

```
curl http://localhost:9090/metrics | grep recommendation_duration
```

## 2. Analyze slow queries

```
SELECT query, mean_exec_time, calls
FROM pg_stat_statements
ORDER BY mean_exec_time DESC
LIMIT 10;
```

## 3. Profile Python code

```
python -m cProfile -s cumulative api_server.py
```

## 4. Check Redis cache hit rate

```
redis-cli INFO stats | grep keyspace
```

### # Solutions

- Add database indexes on frequently queried columns
- Implement query result caching
- Use connection pooling
- Optimize N+1 queries with eager loading
- Add CDN for static assets

## **Issue: Memory leaks in ML model serving**

### # Detection

```
import tracemalloc
import gc
```

```
tracemalloc.start()
```

```
# ... run your application ...
```

```
snapshot = tracemalloc.take_snapshot()
top_stats = snapshot.statistics('lineno')
```

```
print("[ Top 10 memory allocations ]")
for stat in top_stats[:10]:
    print(stat)
```

### # Solution: Proper cleanup

```
class ModelServer:
    def __init__(self):
        self.models = {}

    def load_model(self, model_name: str):
        """Load model with proper memory management"""
```

```

# Clear any existing model
if model_name in self.models:
    del self.models[model_name]
    gc.collect()
    torch.cuda.empty_cache() # If using GPU

# Load new model
self.models[model_name] = load_model(model_name)

def __del__(self):
    """Cleanup on deletion"""
    for model in self.models.values():
        del model
    gc.collect()

```

## 5. Development Resources

### Useful Libraries and Tools

Python:

web\_frameworks:

- FastAPI: Modern async web framework
- Django: Full-featured web framework
- Flask: Lightweight web framework

ml\_libraries:

- PyTorch: Deep learning framework
- scikit-learn: Traditional ML
- XGBoost: Gradient boosting
- Ray: Distributed computing

data\_processing:

- Pandas: Data manipulation
- NumPy: Numerical computing
- Polars: Fast DataFrame library
- Dask: Parallel computing

testing:

- pytest: Testing framework
- hypothesis: Property-based testing
- locust: Load testing
- pytest-asyncio: Async testing

## JavaScript/TypeScript:

### frameworks:

- Next.js: React framework
- Remix: Full-stack framework
- Vite: Build tool

### ui\_libraries:

- Radix UI: Headless components
- Tailwind CSS: Utility CSS
- Framer Motion: Animations
- Recharts: Data visualization

### testing:

- Jest: Testing framework
- React Testing Library: Component testing
- Cypress: E2E testing
- Playwright: Browser automation

## Infrastructure:

### containers:

- Docker: Containerization
- Kubernetes: Orchestration
- Helm: K8s package manager

### monitoring:

- Prometheus: Metrics
- Grafana: Visualization
- Jaeger: Distributed tracing
- ELK Stack: Logging

### databases:

- PostgreSQL: Relational DB
- Neo4j: Graph DB
- Redis: Cache/Message broker
- ClickHouse: Analytics DB

## API Documentation Standards

openapi: 3.0.0

info:

title: EASY-Diabetes Clinical API

version: 2.0.0

description: Clinical decision support API for diabetes management

contact:

email: api-support@easy-diabetes.com  
license:  
name: Proprietary

servers:

- url: https://api.easy-diabetes.com/v2  
description: Production server
- url: https://staging-api.easy-diabetes.com/v2  
description: Staging server

security:

- bearerAuth: []
- oauth2: [read, write]

paths:

/patients/{patientId}/recommendations:

post:

summary: Generate clinical recommendations

operationId: generateRecommendations

tags:

- Clinical Decision Support

parameters:

- name: patientId

in: path

required: true

schema:

type: string

format: uuid

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/RecommendationRequest'

responses:

'200':

description: Recommendations generated successfully

content:

application/json:

schema:

\$ref: '#/components/schemas/RecommendationResponse'

'400':

\$ref: '#/components/responses/BadRequest'

'401':

```
$ref: '#/components/responses/Unauthorized'
'404':
  $ref: '#/components/responses/NotFound'
'500':
  $ref: '#/components/responses/InternalServerError'
```

## Git Workflow

# Feature branch workflow

```
git checkout -b feature/EASY-123-glucose-prediction
```

# Make changes with conventional commits

```
git commit -m "feat(prediction): add 4-hour glucose forecasting"
```

```
git commit -m "test(prediction): add unit tests for glucose predictor"
```

```
git commit -m "docs(prediction): update API documentation"
```

# Keep branch updated

```
git checkout main
```

```
git pull origin main
```

```
git checkout feature/EASY-123-glucose-prediction
```

```
git rebase main
```

# Push and create PR

```
git push origin feature/EASY-123-glucose-prediction
```

# After review and approval

```
git checkout main
```

```
git merge --no-ff feature/EASY-123-glucose-prediction
```

```
git push origin main
```

# Tag releases

```
git tag -a v2.1.0 -m "Release version 2.1.0"
```

```
git push origin v2.1.0
```

## Deployment Checklist

## Pre-Deployment Checklist

### Code Quality

- [ ] All tests passing (unit, integration, e2e)
- [ ] Code coverage >80%
- [ ] No critical security vulnerabilities
- [ ] Performance benchmarks met



- [ ] Documentation updated

#### ### Database

- [ ] Migration scripts tested
- [ ] Rollback plan prepared
- [ ] Indexes optimized
- [ ] Backup completed

#### ### Infrastructure

- [ ] Resource limits configured
- [ ] Autoscaling policies set
- [ ] Health checks defined
- [ ] Monitoring alerts configured
- [ ] SSL certificates valid

#### ### Security

- [ ] Security scan completed
- [ ] Secrets rotated
- [ ] Access controls verified
- [ ] Audit logging enabled
- [ ] WAF rules updated

#### ### Business

- [ ] Stakeholders notified
- [ ] Downtime window scheduled
- [ ] Support team briefed
- [ ] Rollback criteria defined
- [ ] Success metrics identified

---

## Document 7: Development Guidelines and Best Practices

### Code Style Guide

#### Python Style Guide

"""

EASY-Diabetes Python Style Guide

Based on PEP 8 with project-specific conventions

```
"""
```

```
# Imports
# Standard library imports first
import os
import sys
from datetime import datetime
from typing import List, Dict, Optional, Union
```

```
# Third-party imports
import numpy as np
import pandas as pd
from fastapi import Fast
```

```
"""
```

```
EASY-Diabetes Python Style Guide (continued)
Based on PEP 8 with project-specific conventions
"""
```

```
# Imports (continued)
# Standard library imports first
import os
import sys
from datetime import datetime
from typing import List, Dict, Optional, Union
```

```
# Third-party imports
import numpy as np
import pandas as pd
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel, Field
from sqlalchemy.ext.asyncio import AsyncSession
```

```
# Local imports
from app.core.config import settings
from app.models.patient import Patient
from app.services.clinical_engine import ClinicalEngine
```

```
# Class definitions
class ClinicalRecommendation(BaseModel):
    """
    Clinical recommendation model.
```

Attributes:

id: Unique identifier  
patient\_id: Patient UUID  
recommendation\_type: Type of recommendation  
priority: Clinical priority (high, medium, low)  
evidence\_level: Strength of evidence (A, B, C, D)  
confidence\_score: ML model confidence (0-1)

"""

```
id: str = Field(..., description="Unique recommendation ID")
patient_id: str = Field(..., description="Patient UUID")
recommendation_type: str = Field(..., regex="^(medication|lifestyle|monitoring)$")
priority: str = Field(..., regex="^(high|medium|low)$")
evidence_level: str = Field(..., regex="^[A-D]$")
confidence_score: float = Field(..., ge=0, le=1)
```

class Config:

```
    schema_extra = {
        "example": {
            "id": "rec_123456",
            "patient_id": "550e8400-e29b-41d4-a716-446655440000",
            "recommendation_type": "medication",
            "priority": "high",
            "evidence_level": "A",
            "confidence_score": 0.92
        }
    }
```

# Function definitions

```
async def calculate_insulin_dose(
    current_glucose: float,
    target_glucose: float,
    insulin_sensitivity_factor: float,
    active_insulin: float = 0
) -> float:
    """
```

Calculate insulin correction dose.

Args:

current\_glucose: Current blood glucose in mg/dL  
target\_glucose: Target blood glucose in mg/dL  
insulin\_sensitivity\_factor: 1 unit drops BG by X mg/dL  
active\_insulin: Units of insulin still active

Returns:

Recommended insulin dose in units

Raises:

ValueError: If glucose values are out of valid range

Example:

```
>>> dose = await calculate_insulin_dose(
...     current_glucose=180,
...     target_glucose=120,
...     insulin_sensitivity_factor=50,
...     active_insulin=0.5
... )
>>> print(f"Recommended dose: {dose} units")
Recommended dose: 0.7 units
"""

if not 40 <= current_glucose <= 600:
    raise ValueError(f"Current glucose {current_glucose} out of valid range")

if not 70 <= target_glucose <= 180:
    raise ValueError(f"Target glucose {target_glucose} out of valid range")

# Calculate correction dose
glucose_delta = current_glucose - target_glucose
correction_dose = glucose_delta / insulin_sensitivity_factor

# Account for active insulin
net_dose = max(0, correction_dose - active_insulin)

# Round to nearest 0.5 unit
return round(net_dose * 2) / 2

# Constants and configuration
class DiabetesConstants:
    """Central location for diabetes-related constants."""

    # Glucose thresholds (mg/dL)
    HYPOGLYCEMIA_LEVEL_1 = 70
    HYPOGLYCEMIA_LEVEL_2 = 54
    TARGET_RANGE_LOW = 70
    TARGET_RANGE_HIGH = 180
    HYPERGLYCEMIA_LEVEL_1 = 180
    HYPERGLYCEMIA_LEVEL_2 = 250
```

```

# HbA1c targets (%)
GENERAL_A1C_TARGET = 7.0
INTENSIVE_A1C_TARGET = 6.5
RELAXED_A1C_TARGET = 8.0

# Time in range targets (%)
TIR_TARGET_MINIMUM = 70
TBR_TARGET_MAXIMUM = 4
TAR_TARGET_MAXIMUM = 25

# Error handling
class ClinicalError(Exception):
    """Base exception for clinical errors."""
    pass

class ContraindicationError(ClinicalError):
    """Raised when a contraindication is detected."""

    def __init__(self, medication: str, reason: str):
        self.medication = medication
        self.reason = reason
        super().__init__(f"Contraindication for {medication}: {reason}")

class InsufficientDataError(ClinicalError):
    """Raised when insufficient data for clinical decision."""

    def __init__(self, data_type: str, minimum_required: int, actual: int):
        self.data_type = data_type
        self.minimum_required = minimum_required
        self.actual = actual
        super().__init__(
            f"Insufficient {data_type} data: required {minimum_required}, got {actual}"
        )

# Async context managers
class DatabaseTransaction:
    """Async context manager for database transactions."""

    def __init__(self, session: AsyncSession):

```

```

        self.session = session

    async def __aenter__(self):
        await self.session.begin()
        return self.session

    async def __aexit__(self, exc_type, exc_val, exc_tb):
        if exc_type:
            await self.session.rollback()
        else:
            await self.session.commit()
        await self.session.close()

# Type hints for complex types
from typing import TypedDict, Protocol, Literal

class GlucoseReading(TypedDict):
    """Type definition for glucose readings."""
    timestamp: datetime
    value: float
    unit: Literal["mg/dL", "mmol/L"]
    source: Literal["cgm", "fingerstick", "lab"]

class PredictionModel(Protocol):
    """Protocol for prediction models."""

    async def predict(self, data: pd.DataFrame) -> np.ndarray:
        """Generate predictions from input data."""
        ...

    def get_feature_importance(self) -> Dict[str, float]:
        """Return feature importance scores."""
        ...

```

## TypeScript/React Style Guide

```

/**
 * EASY-Diabetes TypeScript/React Style Guide
 * Based on Airbnb style guide with project conventions
 */

```

```

// Import organization

```

```

// 1. React imports
import React, { useState, useEffect, useCallback, useMemo } from 'react'
import { useRouter } from 'next/router'

// 2. Third-party libraries
import { useQuery, useMutation } from '@tanstack/react-query'
import { z } from 'zod'
import { format, parseISO } from 'date-fns'

// 3. Local components
import { Button } from '@components/ui/button'
import { Card, CardContent, CardHeader } from '@components/ui/card'

// 4. Utils and helpers
import { api } from '@lib/api'
import { cn } from '@lib/Utils'

// 5. Types
import type { Patient, ClinicalData, Medication } from '@types/clinical'

// Component definition with proper typing
interface PatientDashboardProps {
  patientId: string
  initialData?: Patient
  onUpdate?: (patient: Patient) => void
}

export const PatientDashboard: React.FC<PatientDashboardProps> = ({
  patientId,
  initialData,
  onUpdate
}) => {
  // State management
  const [isEditing, setIsEditing] = useState(false)
  const [selectedMedication, setSelectedMedication] = useState<Medication | null>(null)

  // Custom hooks
  const router = useRouter()

  // Data fetching with React Query
  const { data: patient, isLoading, error } = useQuery({
    queryKey: ['patient', patientId],
    queryFn: () => api.patients.get(patientId),
  })

```

```

    initialData,
    staleTime: 5 * 60 * 1000, // 5 minutes
  })

  // Mutations
  const updatePatient = useMutation({
    mutationFn: (data: Partial<Patient>) =>
      api.patients.update(patientId, data),
    onSuccess: (updatedPatient) => {
      onUpdate?.(updatedPatient)
      setIsEditing(false)
    },
  })

  // Memoized calculations
  const riskScore = useMemo(() => {
    if (!patient) return null
    return calculateRiskScore(patient)
  }, [patient])

  // Callbacks to prevent re-renders
  const handleEdit = useCallback(() => {
    setIsEditing(true)
  }, [])

  const handleSave = useCallback(async (data: Partial<Patient>) => {
    await updatePatient.mutateAsync(data)
  }, [updatePatient])

  // Effects
  useEffect(() => {
    if (error) {
      console.error('Failed to load patient:', error)
      router.push('/patients')
    }
  }, [error, router])

  // Loading state
  if (isLoading) {
    return <PatientDashboardSkeleton />
  }

  // Error state
  if (error) {

```



```

return (
  <Card>
    <CardContent>
      <p className="text-destructive">
        Failed to load patient data. Please try again.
      </p>
    </CardContent>
  </Card>
)
}

// Main render
return (
  <div className="space-y-6">
    <PatientHeader
      patient={patient!}
      isEditing={isEditing}
      onEdit={handleEdit}
      onSave={handleSave}
    />

    <div className="grid gap-6 md:grid-cols-2 lg:grid-cols-3">
      <GlycemicControlCard patientId={patientId} />
      <MedicationsCard
        medications={patient!.medications}
        onSelect={setSelectedMedication}
      />
      <RiskAssessmentCard score={riskScore} />
    </div>

    {selectedMedication && (
      <MedicationDetailsDialog
        medication={selectedMedication}
        onClose={() => setSelectedMedication(null)}
      />
    )}
  </div>
)
}

// Subcomponents with proper prop types
interface GlycemicControlCardProps {
  patientId: string
}

```

```

const GlycemicControlCard: React.FC<GlycemicControlCardProps> = ({
  patientId
}) => {
  const { data: glucoseData } = useQuery({
    queryKey: ['glucose', patientId, 'recent'],
    queryFn: () => api.glucose.getRecent(patientId),
  })

  return (
    <Card>
      <CardHeader>
        <h3 className="text-lg font-semibold">Glycemic Control</h3>
      </CardHeader>
      <CardContent>
        {glucoseData ? (
          <GlucoseChart data={glucoseData} />
        ) : (
          <div className="h-32 animate-pulse bg-muted rounded" />
        )}
      </CardContent>
    </Card>
  )
}

```

```

// Utility functions with proper typing
function calculateRiskScore(patient: Patient): number {
  let score = 0

  // Age factor
  const age = calculateAge(patient.dateOfBirth)
  if (age > 65) score += 2
  else if (age > 45) score += 1

  // Comorbidity factors
  if (patient.conditions.includes('hypertension')) score += 1
  if (patient.conditions.includes('dyslipidemia')) score += 1

  // Recent HbA1c
  const recentA1c = patient.labs.hba1c[0]?.value
  if (recentA1c > 9) score += 3
  else if (recentA1c > 8) score += 2
  else if (recentA1c > 7) score += 1
}

```

```

    return Math.min(score, 10) // Cap at 10
  }

// Custom hooks
export function useGlucosePrediction(patientId: string) {
  const [prediction, setPrediction] = useState<GlucosePrediction | null>(null)
  const [isLoading, setIsLoading] = useState(false)

  const generatePrediction = useCallback(async () => {
    setIsLoading(true)
    try {
      const result = await api.predictions.generateGlucose(patientId)
      setPrediction(result)
    } catch (error) {
      console.error('Prediction failed:', error)
    } finally {
      setIsLoading(false)
    }
  }, [patientId])

  useEffect(() => {
    generatePrediction()

    // Refresh every 5 minutes
    const interval = setInterval(generatePrediction, 5 * 60 * 1000)
    return () => clearInterval(interval)
  }, [generatePrediction])

  return { prediction, isLoading, refresh: generatePrediction }
}

// Zod schemas for validation
const MedicationFormSchema = z.object({
  name: z.string().min(1, 'Medication name is required'),
  dose: z.number().positive('Dose must be positive'),
  unit: z.enum(['mg', 'g', 'mcg', 'units', 'mL']),
  frequency: z.enum(['daily', 'bid', 'tid', 'qid', 'prn']),
  route: z.enum(['oral', 'subcutaneous', 'intravenous', 'intramuscular']),
  startDate: z.date(),
  endDate: z.date().optional(),
})

type MedicationFormData = z.infer<typeof MedicationFormSchema>

```

```
// Constants
export const GLUCOSE_THRESHOLDS = {
  hypoglycemia: {
    level1: 70,
    level2: 54,
  },
  target: {
    low: 70,
    high: 180,
  },
  hyperglycemia: {
    level1: 180,
    level2: 250,
  },
} as const

// Enums with const assertion
export const MedicationClass = {
  BIGUANIDE: 'biguanide',
  SULFONYLUREA: 'sulfonylurea',
  DPP4_INHIBITOR: 'dpp4_inhibitor',
  SGLT2_INHIBITOR: 'sglt2_inhibitor',
  GLP1_AAGONIST: 'glp1_agonist',
  INSULIN: 'insulin',
} as const

export type MedicationClass = typeof MedicationClass[keyof typeof MedicationClass]
```

## SQL Style Guide

```
-- EASY-Diabetes SQL Style Guide
-- Consistent formatting for maintainable database code

-- Table creation with proper structure
CREATE TABLE clinical.patient_medications (
  -- Primary key first
  id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),

  -- Foreign keys
  patient_id UUID NOT NULL REFERENCES patient.patients(id) ON DELETE CASCADE,
  medication_id UUID NOT NULL REFERENCES clinical.medications(id),
  prescriber_id UUID REFERENCES auth.users(id),

  -- Required fields
```

```

dose NUMERIC(10,3) NOT NULL CHECK (dose > 0),
dose_unit VARCHAR(20) NOT NULL,
frequency VARCHAR(50) NOT NULL,
route VARCHAR(30) NOT NULL,
start_date DATE NOT NULL,

-- Optional fields
end_date DATE,
discontinuation_reason TEXT,
instructions TEXT,

-- Metadata
is_active BOOLEAN DEFAULT true,
created_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMPTZ DEFAULT CURRENT_TIMESTAMP,
created_by UUID REFERENCES auth.users(id),

-- Constraints
CONSTRAINT valid_date_range CHECK (end_date IS NULL OR end_date >= start_date),
CONSTRAINT valid_dose_unit CHECK (dose_unit IN ('mg', 'g', 'mcg', 'units', 'mL'))
);

-- Indexes for performance
CREATE INDEX idx_patient_medications_patient_active
  ON clinical.patient_medications(patient_id, is_active)
  WHERE is_active = true;

CREATE INDEX idx_patient_medications_dates
  ON clinical.patient_medications(patient_id, start_date, end_date);

-- Comments for documentation
COMMENT ON TABLE clinical.patient_medications IS 'Tracks all patient medications with dosing information';
COMMENT ON COLUMN clinical.patient_medications.dose IS 'Medication dose amount';
COMMENT ON COLUMN clinical.patient_medications.frequency IS 'Dosing frequency (e.g., daily, BID, TID)';

-- Complex queries with CTEs for readability
WITH active_medications AS (
  -- Get all active medications for patients
  SELECT
    pm.patient_id,
    pm.medication_id,
    m.generic_name,

```

```

        m.drug_class,
        pm.dose,
        pm.dose_unit,
        pm.frequency,
        pm.start_date,
        DATE_PART('day', CURRENT_DATE - pm.start_date) AS days_on_medication
FROM clinical.patient_medications pm
INNER JOIN clinical.medications m ON m.id = pm.medication_id
WHERE pm.is_active = true
      AND pm.end_date IS NULL
),
medication_counts AS (
  -- Count medications by class
  SELECT
    patient_id,
    drug_class,
    COUNT(*) AS medication_count,
    STRING_AGG(generic_name, ' ' ORDER BY generic_name) AS medications
  FROM active_medications
  GROUP BY patient_id, drug_class
),
patient_summary AS (
  -- Summarize patient medication profile
  SELECT
    p.id AS patient_id,
    p.external_id,
    COUNT(DISTINCT am.medication_id) AS total_medications,
    COUNT(DISTINCT am.drug_class) AS unique_drug_classes,
    MAX(am.days_on_medication) AS longest_medication_duration
  FROM patient.patients p
  LEFT JOIN active_medications am ON am.patient_id = p.id
  WHERE p.is_active = true
  GROUP BY p.id, p.external_id
)
-- Final query combining CTEs
SELECT
  ps.external_id,
  ps.total_medications,
  ps.unique_drug_classes,
  ps.longest_medication_duration,
  mc.drug_class,
  mc.medication_count,
  mc.medications
FROM patient_summary ps

```

```

LEFT JOIN medication_counts mc ON mc.patient_id = ps.patient_id
ORDER BY ps.external_id, mc.drug_class;

-- Function with proper error handling
CREATE OR REPLACE FUNCTION clinical.calculate_medication_adherence(
    p_patient_id UUID,
    p_medication_id UUID,
    p_start_date DATE DEFAULT CURRENT_DATE - INTERVAL '30 days',
    p_end_date DATE DEFAULT CURRENT_DATE
)
RETURNS TABLE (
    adherence_rate NUMERIC,
    expected_doses INTEGER,
    actual_doses INTEGER,
    missed_doses INTEGER
)
LANGUAGE plpgsql
AS $$
DECLARE
    v_frequency VARCHAR(50);
    v_daily_doses INTEGER;
BEGIN
    -- Validate inputs
    IF p_end_date <= p_start_date THEN
        RAISE EXCEPTION 'End date must be after start date';
    END IF;

    -- Get medication frequency
    SELECT pm.frequency INTO v_frequency
    FROM clinical.patient_medications pm
    WHERE pm.patient_id = p_patient_id
        AND pm.medication_id = p_medication_id
        AND pm.is_active = true;

    IF v_frequency IS NULL THEN
        RAISE EXCEPTION 'No active medication found for patient % and medication %',
            p_patient_id, p_medication_id;
    END IF;

    -- Calculate daily doses based on frequency
    v_daily_doses := CASE v_frequency
        WHEN 'daily' THEN 1
        WHEN 'bid' THEN 2
        WHEN 'tid' THEN 3

```

```

        WHEN 'qid' THEN 4
        ELSE 1 -- Default to daily
    END;

-- Calculate adherence
RETURN QUERY
WITH dose_counts AS (
    SELECT
        COUNT(*) AS actual_doses
    FROM clinical.medication_administrations ma
    WHERE ma.patient_id = p_patient_id
        AND ma.medication_id = p_medication_id
        AND ma.administered_at::date BETWEEN p_start_date AND p_end_date
)
SELECT
    ROUND((dc.actual_doses::numeric /
        (v_daily_doses * (p_end_date - p_start_date + 1))::numeric) * 100, 2) AS
adherence_rate,
    v_daily_doses * (p_end_date - p_start_date + 1) AS expected_doses,
    dc.actual_doses::integer AS actual_doses,
    (v_daily_doses * (p_end_date - p_start_date + 1) - dc.actual_doses)::integer AS
missed_doses
    FROM dose_counts dc;
END;
$$;

-- Grant appropriate permissions
GRANT EXECUTE ON FUNCTION clinical.calculate_medication_adherence TO clinical_user;

```

## Testing Strategy

### Unit Testing

```

# tests/unit/test_glucose_prediction.py
import pytest
from datetime import datetime, timedelta
import numpy as np
from unittest.mock import Mock, patch

from app.services.prediction import GlucosePredictor
from app.models.glucose import GlucoseReading
from app.exceptions import InsufficientDataError

```



```

class TestGlucosePredictor:
    """Comprehensive unit tests for glucose prediction."""

    @pytest.fixture
    def predictor(self):
        """Create predictor instance for testing."""
        return GlucosePredictor()

    @pytest.fixture
    def sample_cgm_data(self):
        """Generate sample CGM data for testing."""
        base_time = datetime.now()
        readings = []

        # Generate 24 hours of 5-minute readings
        for i in range(288):
            time = base_time - timedelta(minutes=5*i)
            # Simulate realistic glucose pattern
            base_glucose = 120
            meal_effect = 30 * np.exp(-((i-60)**2)/1000) if 50 < i < 70 else 0
            circadian = 10 * np.sin(2*np.pi*i/288)
            noise = np.random.normal(0, 5)

            glucose = base_glucose + meal_effect + circadian + noise
            glucose = np.clip(glucose, 40, 400)

            readings.append(GlucoseReading(
                timestamp=time,
                value=glucose,
                source='cgm'
            ))

        return readings

    def test_prediction_with_sufficient_data(self, predictor, sample_cgm_data):
        """Test prediction with adequate historical data."""
        # Act
        predictions = predictor.predict(
            historical_data=sample_cgm_data,
            horizon_minutes=60
        )

        # Assert

```

```

assert len(predictions) == 12 # 60 min / 5 min intervals
assert all(40 <= p.value <= 400 for p in predictions)
assert all(p.confidence > 0.7 for p in predictions[:6]) # Higher confidence for near term

def test_prediction_with_insufficient_data(self, predictor):
    """Test that prediction fails with insufficient data."""
    # Arrange
    insufficient_data = [
        GlucoseReading(
            timestamp=datetime.now(),
            value=120,
            source='cgm'
        )
    ]

    # Act & Assert
    with pytest.raises(InsufficientDataError) as exc_info:
        predictor.predict(insufficient_data, horizon_minutes=60)

    assert exc_info.value.minimum_required == 12 # 1 hour of data minimum
    assert exc_info.value.actual == 1

@pytest.mark.parametrize("horizon,expected_confidence", [
    (15, 0.95), # 15 min - very high confidence
    (30, 0.90), # 30 min - high confidence
    (60, 0.80), # 60 min - moderate confidence
    (120, 0.65), # 120 min - lower confidence
    (240, 0.50), # 240 min - uncertainty threshold
])
def test_confidence_decreases_with_horizon(
    self, predictor, sample_cgm_data, horizon, expected_confidence
):
    """Test that prediction confidence decreases with longer horizons."""
    # Act
    predictions = predictor.predict(sample_cgm_data, horizon)

    # Assert
    avg_confidence = np.mean([p.confidence for p in predictions])
    assert abs(avg_confidence - expected_confidence) < 0.1

def test_hypoglycemia_detection(self, predictor):
    """Test that predictor detects impending hypoglycemia."""
    # Arrange - Create declining glucose trend
    base_time = datetime.now()

```

```

declining_data = []

for i in range(24): # 2 hours of data
    time = base_time - timedelta(minutes=5*i)
    glucose = 120 - (3 * i) # Declining by 3 mg/dL every 5 min
    declining_data.append(GlucoseReading(
        timestamp=time,
        value=max(glucose, 40),
        source='cgm'
    ))

# Act
predictions = predictor.predict(declining_data, horizon_minutes=30)
alerts = predictor.detect_hypoglycemia_risk(predictions)

# Assert
assert any(alert.severity == 'high' for alert in alerts)
assert any(p.value < 70 for p in predictions)

@patch('app.services.prediction.ml_model')
def test_model_fallback_on_error(self, mock_model, predictor, sample_cgm_data):
    """Test fallback to statistical model if ML model fails."""
    # Arrange
    mock_model.predict.side_effect = RuntimeError("Model error")

    # Act
    predictions = predictor.predict(sample_cgm_data, horizon_minutes=30)

    # Assert
    assert len(predictions) == 6 # Still returns predictions
    assert all(p.model_type == 'statistical_fallback' for p in predictions)

```

## Integration Testing

```

# tests/integration/test_clinical_workflow.py
import pytest
from httpx import AsyncClient
from sqlalchemy.ext.asyncio import AsyncSession

from app.main import app
from app.models.patient import Patient
from app.models.clinical import ClinicalData
from tests.factories import PatientFactory, ClinicalDataFactory

```

```

@pytest.mark.asyncio
class TestClinicalWorkflow:
    """Integration tests for complete clinical workflows."""

    async def test_complete_recommendation_workflow(
        self,
        async_client: AsyncClient,
        db_session: AsyncSession,
        authenticated_headers: dict
    ):
        """Test end-to-end recommendation generation workflow."""
        # Arrange - Create test patient with clinical data
        patient = await PatientFactory.create(db_session)
        clinical_data = await ClinicalDataFactory.create_batch(
            db_session,
            size=10,
            patient_id=patient.id
        )

        # Act - Generate recommendations
        response = await async_client.post(
            f"/api/v1/patients/{patient.id}/recommendations",
            json={
                "include_categories": ["medication", "lifestyle", "monitoring"],
                "urgency": "routine",
                "context": "regular_checkup"
            },
            headers=authenticated_headers
        )

        # Assert - Check response
        assert response.status_code == 200
        recommendations = response.json()

        assert "recommendations" in recommendations
        assert len(recommendations["recommendations"]) > 0
        assert all(
            rec["evidence_level"] in ["A", "B", "C", "D"]
            for rec in recommendations["recommendations"]
        )
        assert recommendations["risk_assessment"] is not None
        assert recommendations["confidence_scores"]["overall"] > 0.7

```

```

# Verify audit trail created
audit_response = await async_client.get(
    f"/api/v1/audit/patient/{patient.id}",
    headers=authenticated_headers
)
assert audit_response.status_code == 200
audit_entries = audit_response.json()
assert any(
    entry["action"] == "generate_recommendations"
    for entry in audit_entries
)

async def test_medication_optimization_with_contraindications(
    self,
    async_client: AsyncClient,
    db_session: AsyncSession,
    authenticated_headers: dict
):
    """Test medication optimization handles contraindications correctly."""
    # Arrange - Patient with CKD contraindicating metformin
    patient = await PatientFactory.create(
        db_session,
        conditions=[{
            "code": "N18.5",
            "display": "Chronic kidney disease, stage 5"
        }]
    )

    # Current medications including metformin
    await db_session.execute(
        """
        INSERT INTO clinical.patient_medications
        (patient_id, medication_id, dose, dose_unit, frequency, route, start_date)
        VALUES
        (:patient_id, :medication_id, 1000, 'mg', 'bid', 'oral', CURRENT_DATE)
        """,
        {
            "patient_id": patient.id,
            "medication_id": "med_metformin_uuid"
        }
    )
    await db_session.commit()

    # Act - Request medication optimization

```

```

response = await async_client.post(
    f"/api/v1/patients/{patient.id}/medications/optimize",
    headers=authenticated_headers
)

# Assert
assert response.status_code == 200
optimization = response.json()

# Should recommend discontinuing metformin
assert any(
    rec["action"] == "discontinue" and
    rec["medication"] == "metformin" and
    "contraindicated" in rec["reason"].lower()
    for rec in optimization["recommendations"]
)

# Should suggest alternatives
alternatives = [
    rec for rec in optimization["recommendations"]
    if rec["action"] == "initiate"
]
assert len(alternatives) > 0
assert all(
    alt["safety_check"]["renal_safe"]
    for alt in alternatives
)

```

## End-to-End Testing

```

// tests/e2e/clinical-dashboard.spec.ts
import { test, expect } from '@playwright/test'
import { loginAs } from './helpers/auth'
import { createTestPatient } from './helpers/data'

test.describe('Clinical Dashboard E2E', () => {
  test.beforeEach(async ({ page }) => {
    // Login as clinician
    await loginAs(page, 'clinician@test.com', 'testpass123')
  })

  test('complete clinical workflow', async ({ page }) => {
    // Create test patient
    const patient = await createTestPatient({

```

```

    name: 'John Doe',
    mrn: 'TEST-12345',
    diagnosis: 'Type 2 Diabetes',
    hba1c: 8.5
  })

  // Navigate to patient dashboard
  await page.goto(`/patients/${patient.id}`)

  // Verify patient information displayed
  await expect(page.locator('h1')).toContainText('John Doe')
  await expect(page.locator('[data-testid="mrn"]')).toContainText('TEST-12345')

  // Check glycemic control section
  await page.click('text=Glycemic Control')
  await expect(page.locator('[data-testid="current-hba1c"]')).toContainText('8.5%')

  // Generate recommendations
  await page.click('button:has-text("Generate Recommendations")')

  // Wait for recommendations to load
  await expect(page.locator('[data-testid="recommendations-loading"]')).toBeVisible()
  await expect(page.locator('[data-testid="recommendations-list"]')).toBeVisible({
    timeout: 10000
  })

  // Verify recommendations displayed
  const recommendations = page.locator('[data-testid="recommendation-card"]')
  await expect(recommendations).toHaveCount(3, { timeout: 5000 })

  // Test accepting a recommendation
  await recommendations.first().locator('button:has-text("Accept")').click()

  // Verify confirmation dialog
  await expect(page.locator('[role="dialog"]')).toContainText('Confirm Recommendation')
  await page.click('button:has-text("Confirm")')

  // Verify success message
  await expect(page.locator('[data-testid="toast"]')).toContainText('Recommendation accepted')

  // Check that recommendation is marked as accepted
  await expect(recommendations.first()).toHaveAttribute('data-status', 'accepted')
})

```

```

test('real-time glucose monitoring', async ({ page }) => {
  const patient = await createTestPatient({
    hasCGM: true,
    cgmDevice: 'Dexcom G7'
  })

  await page.goto(`/patients/${patient.id}/glucose`)

  // Verify CGM data is displayed
  await expect(page.locator('[data-testid="cgm-chart"]')).toBeVisible()
  await expect(page.locator('[data-testid="current-glucose"]')).toContainText(/d+ mg/dL/)

  // Test prediction feature
  await page.click('button:has-text("Show Prediction")')
  await expect(page.locator('[data-testid="glucose-prediction"]')).toBeVisible()

  // Verify prediction confidence
  const confidenceText = await
page.locator('[data-testid="prediction-confidence"]').textContent()
  const confidence = parseFloat(confidenceText?.match(/(\d+)\%/)?.[1] || '0')
  expect(confidence).toBeGreaterThan(70)

  // Test alert settings
  await page.click('button:has-text("Alert Settings")')
  await page.fill('#low-alert-threshold', '75')
  await page.fill('#high-alert-threshold', '170')
  await page.click('button:has-text("Save Settings")')

  await expect(page.locator('[data-testid="toast"]')).toContainText('Alert settings updated')
})
})

```

## Deployment and Operations

### Docker Configuration

```

# Base image for Python services
FROM python:3.11-slim as python-base

# Install system dependencies
RUN apt-get update && apt-get install -y \
  build-essential \
  curl \

```



```

libpq-dev \
&& rm -rf /var/lib/apt/lists/*

# Set environment variables
ENV PYTHONUNBUFFERED=1 \
    PYTHONDONTWRITEBYTECODE=1 \
    PIP_NO_CACHE_DIR=1 \
    PIP_DISABLE_PIP_VERSION_CHECK=1

# Install Python dependencies
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Production image for clinical API
FROM python-base as clinical-api

# Copy application code
COPY ./services/clinical-api /app

# Create non-root user
RUN useradd -m -u 1000 appuser && chown -R appuser:appuser /app
USER appuser

# Health check
HEALTHCHECK --interval=30s --timeout=3s --retries=3 \
    CMD curl -f http://localhost:8000/health || exit 1

# Run application
CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]

# ML model serving image
FROM python-base as model-server

# Install ML-specific dependencies
RUN pip install --no-cache-dir \
    torch==2.1.0 \
    transformers==4.35.0 \
    onnxruntime==1.16.0

# Copy model files and code
COPY ./services/model-server /app
COPY ./models /models

```

```
# Create non-root user
RUN useradd -m -u 1000 appuser && chown -R appuser:appuser /app /models
USER appuser

# Configure model cache
ENV TRANSFORMERS_CACHE=/models/cache \
    TORCH_HOME=/models/torch

# Run model server
CMD ["python", "-m", "app.server", "--port", "8001"]
```

## Kubernetes Manifests

```
# kubernetes/base/deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: clinical-api
  labels:
    app: clinical-api
    version: v2.0.0
spec:
  replicas: 3
  selector:
    matchLabels:
      app: clinical-api
  template:
    metadata:
      labels:
        app: clinical-api
        version: v2.0.0
    annotations:
      prometheus.io/scrape: "true"
      prometheus.io/port: "8000"
      prometheus.io/path: "/metrics"
    spec:
      serviceAccountName: clinical-api
      securityContext:
        runAsNonRoot: true
        runAsUser: 1000
        fsGroup: 1000
      containers:
        - name: clinical-api
          image: easy-diabetes/clinical-api:v2.0.0
```

```
imagePullPolicy: IfNotPresent
ports:
- name: http
  containerPort: 8000
  protocol: TCP
env:
- name: ENV
  value: "production"
- name: DATABASE_URL
  valueFrom:
    secretKeyRef:
      name: clinical-api-secrets
      key: database-url
- name: REDIS_URL
  valueFrom:
    secretKeyRef:
      name: clinical-api-secrets
      key: redis-url
resources:
  requests:
    memory: "512Mi"
    cpu: "250m"
  limits:
    memory: "1Gi"
    cpu: "1000m"
livenessProbe:
  httpGet:
    path: /health
    port: http
  initialDelaySeconds: 30
  periodSeconds: 10
  timeoutSeconds: 3
  failureThreshold: 3
readinessProbe:
  httpGet:
    path: /ready
    port: http
  initialDelaySeconds: 5
  periodSeconds: 5
  timeoutSeconds: 3
  successThreshold: 1
volumeMounts:
- name: config
  mountPath: /app/config
```

```

      readOnly: true
    volumes:
      - name: config
        configMap:
          name: clinical-api-config
---
apiVersion: v1
kind: Service
metadata:
  name: clinical-api
  labels:
    app: clinical-api
spec:
  type: ClusterIP
  ports:
    - port: 80
      targetPort: http
      protocol: TCP
      name: http
  selector:
    app: clinical-api
---
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: clinical-api-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: clinical-api
  minReplicas: 3
  maxReplicas: 20
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70
    - type: Resource
      resource:
        name: memory
        target:

```

```
    type: Utilization
    averageUtilization: 80
- type: Pods
  pods:
    metric:
      name: http_requests_per_second
    target:
      type: AverageValue
      averageValue: "1000"
behavior:
  scaleDown:
    stabilizationWindowSeconds: 300
    policies:
      - type: Percent
        value: 10
        periodSeconds: 60
  scaleUp:
    stabilizationWindowSeconds: 60
    policies:
      - type: Percent
        value: 50
        periodSeconds: 60
      - type: Pods
        value: 2
        periodSeconds: 60
```

## CI/CD Pipeline

```
# .github/workflows/deploy.yml
name: Deploy to Production
```

```
on:
  push:
    branches: [main]
  workflow_dispatch:
```

```
env:
  REGISTRY: ghcr.io
  IMAGE_NAME: ${GITHUB_REPOSITORY}
```

```
jobs:
  test:
    runs-on: ubuntu-latest
    services:
```

postgres:  
image: postgres:16  
env:  
 POSTGRES\_PASSWORD: testpass  
options: >-  
 --health-cmd pg\_isready  
 --health-interval 10s  
 --health-timeout 5s  
 --health-retries 5

redis:  
image: redis:7-alpine  
options: >-  
 --health-cmd "redis-cli ping"  
 --health-interval 10s  
 --health-timeout 5s  
 --health-retries 5

steps:

- uses: actions/checkout@v4

- name: Set up Python

uses: actions/setup-python@v4

with:

python-version: '3.11'

cache: 'pip'

- name: Install dependencies

run: |

python -m pip install --upgrade pip

pip install -r requirements-test.txt

- name: Run tests

env:

DATABASE\_URL: postgresql://postgres:testpass@localhost/test\_db

REDIS\_URL: redis://localhost:6379

run: |

pytest -v --cov=app --cov-report=xml

- name: Upload coverage

uses: codecov/codecov-action@v3

with:

file: ./coverage.xml

build:

```

needs: test
runs-on: ubuntu-latest
permissions:
  contents: read
  packages: write

steps:
- uses: actions/checkout@v4

- name: Set up Docker Buildx
  uses: docker/setup-buildx-action@v3

- name: Log in to Container Registry
  uses: docker/login-action@v3
  with:
    registry: ${{ env.REGISTRY }}
    username: ${{ github.actor }}
    password: ${{ secrets.GITHUB_TOKEN }}

- name: Extract metadata
  id: meta
  uses: docker/metadata-action@v5
  with:
    images: ${{ env.REGISTRY }}/${{ env.IMAGE_NAME }}
    tags: |
      type=ref,event=branch
      type=ref,event=pr
      type=semver,pattern={{version}}
      type=semver,pattern={{major}}.{{minor}}
      type=sha,prefix={{branch}}-

- name: Build and push Docker image
  uses: docker/build-push-action@v5
  with:
    context: .
    platforms: linux/amd64,linux/arm64
    push: true
    tags: ${{ steps.meta.outputs.tags }}
    labels: ${{ steps.meta.outputs.labels }}
    cache-from: type=gha
    cache-to: type=gha,mode=max

deploy:
  needs: build

```

runs-on: ubuntu-latest  
environment: production

steps:

- uses: actions/checkout@v4

- name: Install kubectl  
uses: azure/setup-kubectl@v3  
with:  
version: 'v1.28.0'

- name: Configure AWS credentials  
uses: aws-actions/configure-aws-credentials@v4  
with:  
aws-access-key-id: \${ secrets.AWS\_ACCESS\_KEY\_ID }  
aws-secret-access-key: \${ secrets.AWS\_SECRET\_ACCESS\_KEY }  
aws-region: us-west-2

- name: Update kubeconfig  
run: |  
aws eks update-kubeconfig --name easy-diabetes-prod --region us-west-2

- name: Deploy to Kubernetes  
run: |  
kubectl set image deployment/clinical-api \  
clinical-api=\${ env.REGISTRY }/\${ env.IMAGE\_NAME }:\${ github.sha } \  
-n production

kubectl rollout status deployment/clinical-api -n production

- name: Run smoke tests  
run: |  
./scripts/smoke-tests.sh https://api.easy-diabetes.com

## Monitoring and Alerting

# prometheus/alerts.yml

groups:

- name: clinical\_api\_alerts

interval: 30s

rules:

- alert: HighErrorRate

expr: |

rate(http\_requests\_total{status=~"5.."}[5m]) > 0.05



```

for: 5m
labels:
  severity: critical
  service: clinical-api
annotations:
  summary: "High error rate detected"
  description: "Error rate is {{ $value | humanizePercentage }} for {{ $labels.instance }}"

- alert: HighLatency
  expr: |
    histogram_quantile(0.95, rate(http_request_duration_seconds_bucket[5m])) > 1
  for: 5m
  labels:
    severity: warning
    service: clinical-api
  annotations:
    summary: "High API latency detected"
    description: "95th percentile latency is {{ $value }}s for {{ $labels.instance }}"

- alert: PredictionServiceDown
  expr: |
    up{job="prediction-service"} == 0
  for: 1m
  labels:
    severity: critical
    service: prediction-service
  annotations:
    summary: "Prediction service is down"
    description: "Prediction service instance {{ $labels.instance }} is down"

- alert: DatabaseConnectionPoolExhausted
  expr: |
    postgresql_connections_active / postgresql_connections_max > 0.9
  for: 5m
  labels:
    severity: warning
    service: database
  annotations:
    summary: "Database connection pool nearly exhausted"
    description: "{{ $value | humanizePercentage }} of connections are in use"

- alert: HighMemoryUsage
  expr: |
    container_memory_usage_bytes / container_spec_memory_limit_bytes > 0.9

```

```

    for: 5m
    labels:
      severity: warning
    annotations:
      summary: "Container memory usage is high"
      description: "Memory usage is {{ $value | humanizePercentage }} for {{ $labels.pod }}"

- name: clinical_sla_alerts
  interval: 1m
  rules:
    - alert: RecommendationGenerationSlow
      expr: |
        histogram_quantile(0.99,
rate(recommendation_generation_duration_seconds_bucket[5m])) > 5
      for: 10m
      labels:
        severity: warning
        sla: performance
      annotations:
        summary: "Recommendation generation is slow"
        description: "99th percentile generation time is {{ $value }}s"

- alert: CriticalDataMissing
  expr: |
    increase(clinical_data_missing_total[1h]) > 10
  for: 5m
  labels:
    severity: critical
    sla: data_quality
  annotations:
    summary: "Critical clinical data missing"
    description: "{{ $value }} instances of missing critical data in the last hour"

```

## Disaster Recovery Plan

# EASY-Diabetes Disaster Recovery Plan

## Recovery Time Objectives (RTO) and Recovery Point Objectives (RPO)

Service	RTO	RPO	Priority
Clinical API	15 min	5 min	P0
Database (Primary)	30 min	1 min	P0
Prediction Service	1 hour	1 hour	P1

| Frontend Application | 15 min | N/A | P0 |  
| Monitoring System | 2 hours | 1 hour | P2 |

## ## Backup Strategy

### ### Database Backups

```
``bash
#!/bin/bash
# Automated backup script (runs every hour)

# Variables
BACKUP_DIR="/backups/postgres"
DB_NAME="easy_diabetes"
TIMESTAMP=$(date +%Y%m%d_%H%M%S)
S3_BUCKET="easy-diabetes-backups"

# Create backup
pg_dump -h $DB_HOST -U $DB_USER -d $DB_NAME -F c -b -v -f
"$BACKUP_DIR/backup_${TIMESTAMP}.dump"

# Compress backup
gzip "$BACKUP_DIR/backup_${TIMESTAMP}.dump"

# Upload to S3
aws s3 cp "$BACKUP_DIR/backup_${TIMESTAMP}.dump.gz" "s3://$S3_BUCKET/postgres/"

# Clean up old local backups (keep 7 days)
find $BACKUP_DIR -name "backup_*.dump.gz" -mtime +7 -delete

# Verify backup integrity
pg_restore --list "$BACKUP_DIR/backup_${TIMESTAMP}.dump.gz" > /dev/null
if [ $? -eq 0 ]; then
    echo "Backup verified successfully"
else
    echo "Backup verification failed" >&2
    exit 1
fi
```

## Application State Backup

```
# Kubernetes CronJob for state backup
apiVersion: batch/v1
kind: CronJob
metadata:
```

```

name: state-backup
spec:
  schedule: "0 */6 * * *" # Every 6 hours
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: backup
              image: easy-diabetes/backup-tool:latest
              command:
                - /bin/bash
                - -c
                - |
                  # Backup Redis state
                  redis-cli --rdb /backup/redis-dump.rdb

                  # Backup ML model cache
                  tar -czf /backup/model-cache.tar.gz /models/cache/

                  # Backup configuration
                  kubectl get configmaps -n production -o yaml > /backup/configmaps.yaml
                  kubectl get secrets -n production -o yaml > /backup/secrets.yaml

                  # Upload to S3
                  aws s3 sync /backup/ s3://easy-diabetes-backups/state/$(date +%Y%m%d)/

```

## Recovery Procedures

### 1. Database Recovery

```

# Restore from backup
BACKUP_FILE="backup_20240115_120000.dump.gz"

# Download from S3
aws s3 cp s3://easy-diabetes-backups/postgres/$BACKUP_FILE /tmp/

# Decompress
gunzip /tmp/$BACKUP_FILE

# Restore database
pg_restore -h $DB_HOST -U $DB_USER -d $DB_NAME -v /tmp/${BACKUP_FILE%.gz}

```

```
# Verify restoration
psql -h $DB_HOST -U $DB_USER -d $DB_NAME -c "SELECT COUNT(*) FROM
patient.patients;"
```

## 2. Application Recovery

```
# Scale down current deployment
kubectl scale deployment clinical-api --replicas=0 -n production
```

```
# Update deployment with recovery image
kubectl set image deployment/clinical-api \
  clinical-api=easy-diabetes/clinical-api:recovery-$VERSION \
  -n production
```

```
# Scale up with new image
kubectl scale deployment clinical-api --replicas=3 -n production
```

```
# Monitor rollout
kubectl rollout status deployment/clinical-api -n production
```

## 3. Multi-Region Failover

```
# Primary region failure detected
# Switch traffic to secondary region
```

```
# Update DNS to point to secondary load balancer
aws route53 change-resource-record-sets \
  --hosted-zone-id $ZONE_ID \
  --change-batch '{
    "Changes": [{
      "Action": "UPSERT",
      "ResourceRecordSet": {
        "Name": "api.easy-diabetes.com",
        "Type": "A",
        "AliasTarget": {
          "HostedZoneId": "Z2FDTNDATAQYW2",
          "DNSName": "secondary-lb.us-east-1.elb.amazonaws.com",
          "EvaluateTargetHealth": true
        }
      }
    }
  ]
}'
```

```
# Verify failover
curl -f https://api.easy-diabetes.com/health
```

## Communication Plan

### Incident Response Team

1. **Incident Commander:** Overall coordination
2. **Technical Lead:** Technical recovery execution
3. **Communications Lead:** Stakeholder updates
4. **Customer Success Lead:** User communication

### Communication Channels

- **Internal:** Slack #incident-response
- **Status Page:** status.easy-diabetes.com
- **Customer Email:** Via automated system
- **Phone Tree:** For P0 incidents

### Status Update Template

**\*\*Incident Status Update\*\***

**\*\*Time\*\*:** [TIMESTAMP]

**\*\*Severity\*\*:** P0/P1/P2

**\*\*Impact\*\*:** [Affected services and users]

**\*\*Current Status\*\*:** Investigating/Identified/Monitoring/Resolved

**\*\*Next Update\*\*:** [TIME]

**\*\*Details\*\*:**

[Brief description of the issue and current actions]

**\*\*Customer Impact\*\*:**

[What users are experiencing]

**\*\*Workaround\*\*** (if available):

[Temporary solutions for users]

**## Security Considerations**

**### Security Architecture**

**``yaml**

```

# Security policies and configurations
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: production
spec:
  mtls:
    mode: STRICT
---
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: clinical-api-authz
  namespace: production
spec:
  selector:
    matchLabels:
      app: clinical-api
  action: ALLOW
  rules:
    - from:
        - source:
            principals: ["cluster.local/ns/production/sa/frontend"]
          to:
            - operation:
                methods: ["GET", "POST"]
                paths: ["/api/v1/*"]
      - from:
          - source:
              principals: ["cluster.local/ns/production/sa/admin"]
            to:
              - operation:
                  methods: ["*"]
---
# Network policies
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: clinical-api-network-policy
  namespace: production
spec:
  podSelector:
    matchLabels:

```

```

    app: clinical-api
policyTypes:
- Ingress
- Egress
ingress:
- from:
  - namespaceSelector:
      matchLabels:
        name: production
  - podSelector:
      matchLabels:
        app: frontend
ports:
- protocol: TCP
  port: 8000
egress:
- to:
  - namespaceSelector:
      matchLabels:
        name: production
ports:
- protocol: TCP
  port: 5432 # PostgreSQL
- protocol: TCP
  port: 6379 # Redis
- to:
  - namespaceSelector: {}
  podSelector:
    matchLabels:
      k8s-app: kube-dns
ports:
- protocol: UDP
  port: 53

```

## HIPAA Compliance Checklist

### # HIPAA Compliance Implementation Checklist

#### ## Administrative Safeguards

#### ### Security Officer and Workforce Training

- [x] Designated HIPAA Security Officer
- [x] Workforce training program implemented
- [x] Access management procedures documented



- [x] Sanction policy for violations
- [x] Periodic security updates training

#### ### Access Management

- [x] Unique user identification
- [x] Automatic logoff (15 min idle)
- [x] Encryption and decryption procedures

### ## Physical Safeguards

#### ### Facility Access Controls

- [x] Data center physical security (AWS/Azure compliance)
- [x] Workstation use policies
- [x] Device and media controls

### ## Technical Safeguards

#### ### Access Control

```
```python
```

```
# Implementation example
from functools import wraps
from typing import Callable
import audit_log
```

```
def hipaa_access_control(resource_type: str, action: str):
    """Decorator for HIPAA-compliant access control."""
    def decorator(func: Callable) -> Callable:
        @wraps(func)
        async def wrapper(*args, **kwargs):
            # Extract user and patient info
            user = kwargs.get('current_user')
            patient_id = kwargs.get('patient_id')

            # Check access permissions
            if not check_access_permission(user, resource_type, action, patient_id):
                # Log access denial
                await audit_log.log_access_denial(
                    user_id=user.id,
                    resource_type=resource_type,
                    action=action,
                    patient_id=patient_id
                )
                raise HTTPException(403, "Access denied")
```

```

        # Log successful access
        await audit_log.log_access(
            user_id=user.id,
            resource_type=resource_type,
            action=action,
            patient_id=patient_id
        )

    # Execute function
    result = await func(*args, **kwargs)

    return result

    return wrapper
return decorator

# Usage
@hipaa_access_control("patient_record", "read")
async def get_patient_record(patient_id: str, current_user: User):
    # Implementation
    pass

```

## Audit Controls

```

-- Audit log table with required HIPAA fields
CREATE TABLE audit.hipaa_audit_log (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    timestamp TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
    user_id UUID NOT NULL,
    user_name VARCHAR(255) NOT NULL,
    user_role VARCHAR(50) NOT NULL,
    action VARCHAR(50) NOT NULL, -- CREATE, READ, UPDATE, DELETE
    resource_type VARCHAR(50) NOT NULL,
    resource_id UUID,
    patient_id UUID,
    phi_accessed TEXT[], -- Specific PHI fields accessed
    access_reason VARCHAR(255),
    ip_address INET NOT NULL,
    user_agent TEXT,
    session_id UUID,
    success BOOLEAN NOT NULL DEFAULT true,
    failure_reason TEXT,

    -- Indexes for audit queries

```

```

    INDEX idx_audit_user_time (user_id, timestamp DESC),
    INDEX idx_audit_patient_time (patient_id, timestamp DESC),
    INDEX idx_audit_action_time (action, timestamp DESC)
) PARTITION BY RANGE (timestamp);

-- Create monthly partitions for audit logs
CREATE TABLE audit.hipaa_audit_log_2024_01
PARTITION OF audit.hipaa_audit_log
FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');

-- Audit report views
CREATE VIEW audit.user_access_report AS
SELECT
    user_id,
    user_name,
    COUNT(*) as total_accesses,
    COUNT(DISTINCT patient_id) as patients_accessed,
    COUNT(DISTINCT DATE(timestamp)) as days_active,
    array_agg(DISTINCT action) as actions_performed
FROM audit.hipaa_audit_log
WHERE timestamp >= CURRENT_DATE - INTERVAL '30 days'
GROUP BY user_id, user_name;

```

## Integrity Controls

```

# Data integrity verification
import hashlib
import hmac
from cryptography.fernet import Fernet

class DataIntegrityService:
    def __init__(self, secret_key: bytes):
        self.secret_key = secret_key
        self.fernet = Fernet(Fernet.generate_key())

    def create_integrity_hash(self, data: dict) -> str:
        """Create HMAC hash for data integrity."""
        # Serialize data consistently
        data_str = json.dumps(data, sort_keys=True)

        # Create HMAC
        h = hmac.new(
            self.secret_key,
            data_str.encode(),

```

```

        hashlib.sha256
    )

    return h.hexdigest()

def verify_integrity(self, data: dict, expected_hash: str) -> bool:
    """Verify data integrity using HMAC."""
    actual_hash = self.create_integrity_hash(data)
    return hmac.compare_digest(actual_hash, expected_hash)

def encrypt_phi(self, phi_data: str) -> str:
    """Encrypt PHI data at rest."""
    return self.fernet.encrypt(phi_data.encode()).decode()

def decrypt_phi(self, encrypted_data: str) -> str:
    """Decrypt PHI data."""
    return self.fernet.decrypt(encrypted_data.encode()).decode()

```

## Transmission Security

```

# TLS configuration for all services
from ssl import create_default_context, Purpose
import certifi

def create_secure_ssl_context():
    """Create SSL context with strict security settings."""
    context = create_default_context(
        purpose=Purpose.CLIENT_AUTH,
        cafile=certifi.where()
    )

    # Force TLS 1.2 minimum
    context.minimum_version = ssl.TLSVersion.TLSv1_2

    # Disable weak ciphers
    context.set_ciphers(
        'ECDHE+AESGCM:ECDHE+CHACHA20:DHE+AESGCM:'
        'DHE+CHACHA20:!aNULL:!MD5:!DSS'
    )

    return context

```

# Performance Optimization

## Database Query Optimization

```
-- Optimized query for patient clinical summary
WITH latest_labs AS (
  SELECT DISTINCT ON (patient_id, lab_type)
    patient_id,
    lab_type,
    value,
    value_numeric,
    unit,
    collected_at
  FROM clinical.lab_results
  WHERE patient_id = ANY($1::uuid[]) -- Array of patient IDs
    AND collected_at >= CURRENT_DATE - INTERVAL '6 months'
  ORDER BY patient_id, lab_type, collected_at DESC
),
active_medications AS (
  SELECT
    pm.patient_id,
    json_agg(
      json_build_object(
        'medication_name', m.generic_name,
        'drug_class', m.drug_class,
        'dose', pm.dose,
        'dose_unit', pm.dose_unit,
        'frequency', pm.frequency,
        'start_date', pm.start_date
      ) ORDER BY pm.start_date DESC
    ) as medications
  FROM clinical.patient_medications pm
  JOIN clinical.medications m ON m.id = pm.medication_id
  WHERE pm.patient_id = ANY($1::uuid[])
    AND pm.is_active = true
  GROUP BY pm.patient_id
),
recent_glucose AS (
  SELECT
    patient_id,
    percentile_cont(0.5) WITHIN GROUP (ORDER BY value) as median_glucose,
    AVG(value) as mean_glucose,
    STDDEV(value) as glucose_variability,
    COUNT(*) as reading_count
```

```

FROM clinical.measurements
WHERE patient_id = ANY($1::uuid[])
  AND measurement_type = 'glucose'
  AND measured_at >= CURRENT_DATE - INTERVAL '14 days'
GROUP BY patient_id
)
SELECT
  p.id,
  p.external_id,
  p.first_name,
  p.last_name,

  -- Latest labs as JSON
  coalesce(
    json_object_agg(
      ll.lab_type,
      json_build_object(
        'value', ll.value,
        'numeric', ll.value_numeric,
        'unit', ll.unit,
        'date', ll.collected_at
      )
    ) FILTER (WHERE ll.lab_type IS NOT NULL),
    '{}'::json
  ) as latest_labs,

  -- Medications
  am.medications,

  -- Recent glucose stats
  json_build_object(
    'median', rg.median_glucose,
    'mean', rg.mean_glucose,
    'variability', rg.glucose_variability,
    'reading_count', rg.reading_count
  ) as glucose_stats

FROM patient.patients p
LEFT JOIN latest_labs ll ON ll.patient_id = p.id
LEFT JOIN active_medications am ON am.patient_id = p.id
LEFT JOIN recent_glucose rg ON rg.patient_id = p.id
WHERE p.id = ANY($1::uuid[])
GROUP BY p.id, p.external_id, p.first_name, p.last_name,
  am.medications, rg.median_glucose, rg.mean_glucose,

```

```

rg.glucose_variability, rg.reading_count;

-- Create covering index for this query
CREATE INDEX idx_clinical_summary ON clinical.lab_results
(patient_id, lab_type, collected_at DESC)
INCLUDE (value, value_numeric, unit)
WHERE collected_at >= CURRENT_DATE - INTERVAL '6 months';

```

## Caching Strategy

```

# Multi-level caching implementation
from typing import Optional, Any, Callable
import asyncio
from functools import wraps
import redis.asyncio as redis
import pickle
import hashlib

class MultiLevelCache:
    def __init__(self):
        self.memory_cache = {} # L1: In-memory cache
        self.redis_client = None # L2: Redis cache
        self.cache_ttl = {
            'short': 300, # 5 minutes
            'medium': 3600, # 1 hour
            'long': 86400, # 24 hours
        }

    async def initialize(self, redis_url: str):
        """Initialize Redis connection."""
        self.redis_client = await redis.from_url(redis_url)

    def cache_key(self, prefix: str, *args, **kwargs) -> str:
        """Generate consistent cache key."""
        key_data = f"{prefix}:{args}:{sorted(kwargs.items())}"
        return hashlib.md5(key_data.encode()).hexdigest()

    async def get(self, key: str) -> Optional[Any]:
        """Get from cache (memory first, then Redis)."""
        # Check L1 cache
        if key in self.memory_cache:
            return self.memory_cache[key]['value']

        # Check L2 cache

```

```

if self.redis_client:
    cached = await self.redis_client.get(key)
    if cached:
        value = pickle.loads(cached)
        # Promote to L1
        self.memory_cache[key] = {'value': value}
        return value

return None

async def set(self, key: str, value: Any, ttl: str = 'medium'):
    """Set in both cache levels."""
    # Set in L1
    self.memory_cache[key] = {'value': value}

    # Set in L2
    if self.redis_client:
        await self.redis_client.setex(
            key,
            self.cache_ttl[ttl],
            pickle.dumps(value)
        )

def cached(self, prefix: str, ttl: str = 'medium'):
    """Decorator for caching function results."""
    def decorator(func: Callable) -> Callable:
        @wraps(func)
        async def wrapper(*args, **kwargs):
            # Generate cache key
            cache_key = self.cache_key(prefix, *args, **kwargs)

            # Check cache
            cached_value = await self.get(cache_key)
            if cached_value is not None:
                return cached_value

            # Call function
            result = await func(*args, **kwargs)

            # Cache result
            await self.set(cache_key, result, ttl)

            return result

```



```
    return wrapper
return decorator
```

```
# Usage example
cache = MultiLevelCache()

@cache.cached("patient_summary", ttl="short")
async def get_patient_summary(patient_id: str) -> dict:
    # Expensive operation
    return await fetch_patient_data(patient_id)
```

## Performance Monitoring

```
# Application performance monitoring
from prometheus_client import Counter, Histogram, Gauge
import time
from contextlib import asynccontextmanager
import structlog
```

```
logger = structlog.get_logger()
```

```
# Metrics
request_count = Counter(
    'http_requests_total',
    'Total HTTP requests',
    ['method', 'endpoint', 'status']
)
```

```
request_duration = Histogram(
    'http_request_duration_seconds',
    'HTTP request duration',
    ['method', 'endpoint']
)
```

```
active_requests = Gauge(
    'http_requests_active',
    'Active HTTP requests'
)
```

```
db_query_duration = Histogram(
    'db_query_duration_seconds',
    'Database query duration',
    ['query_type', 'table']
)
```

```

cache_operations = Counter(
    'cache_operations_total',
    'Cache operations',
    ['operation', 'result']
)

@asynccontextmanager
async def track_performance(operation: str, **tags):
    """Context manager for tracking operation performance."""
    start_time = time.time()

    # Log start
    logger.info(f"{operation}_started", **tags)

    try:
        yield

        # Record success
        duration = time.time() - start_time
        logger.info(
            f"{operation}_completed",
            duration=duration,
            **tags
        )

    except Exception as e:
        # Record failure
        duration = time.time() - start_time
        logger.error(
            f"{operation}_failed",
            duration=duration,
            error=str(e),
            **tags
        )
        raise

# Middleware for request tracking
async def performance_middleware(request, call_next):
    """Track request performance."""
    active_requests.inc()
    start_time = time.time()

    try:

```

```

response = await call_next(request)
duration = time.time() - start_time

# Record metrics
request_count.labels(
    method=request.method,
    endpoint=request.url.path,
    status=response.status_code
).inc()

request_duration.labels(
    method=request.method,
    endpoint=request.url.path
).observe(duration)

# Add performance headers
response.headers["X-Response-Time"] = f"{duration:.3f}"

return response

finally:
    active_requests.dec()

```

## Conclusion

This comprehensive documentation suite provides the foundation for developing EASY-Diabetes, a state-of-the-art clinical decision support system. The documentation covers:

1. **System Architecture:** Microservices design with scalable, fault-tolerant infrastructure
2. **Medical Knowledge Base:** Comprehensive clinical guidelines and evidence-based protocols
3. **Technical Implementation:** Detailed code examples and best practices
4. **Security and Compliance:** HIPAA-compliant security measures and audit trails
5. **Testing Strategy:** Unit, integration, and end-to-end testing approaches
6. **Deployment and Operations:** CI/CD pipelines and monitoring strategies
7. **Performance Optimization:** Caching, query optimization, and monitoring

The system is designed to handle millions of users while maintaining sub-second response times and 99.9% availability. The modular architecture allows for easy extension to other chronic conditions beyond diabetes.

**Document Version:** 2.0.0

**Last Updated:** January 2025

**Next Review:** January 2026

**Maintainers:** EASY-Diabetes Development Team